

# Co-Designing Distributed Systems with Networks

CS5229 Guest Lecture

Jialin Li



**NUS** | Computing

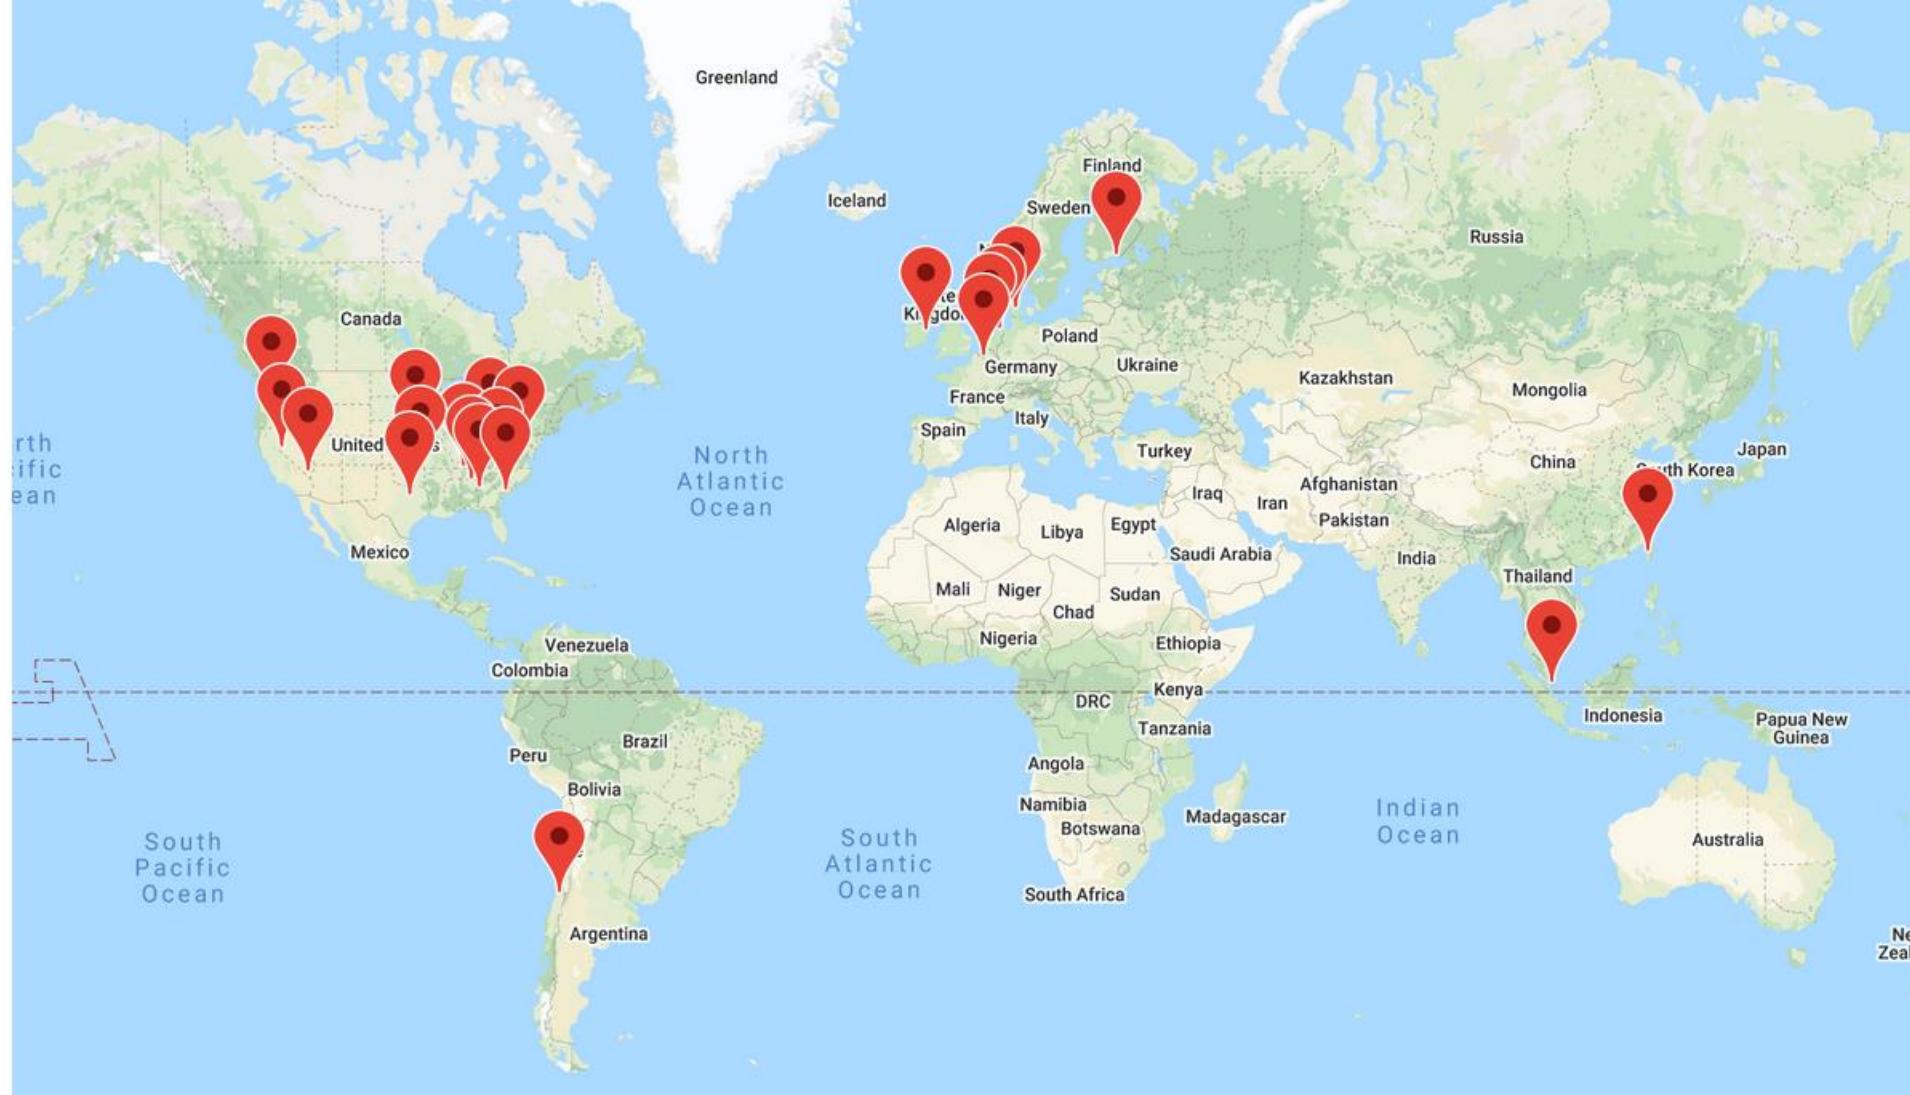
National University  
of Singapore

# About me

- B.S.E at the University of Michigan
- Ph.D. at the University of Washington
- Started faculty job at NUS in 2020
- Research interests: distributed systems, operating systems

# Distributed Systems are

E



# Building Distributed Systems is Challenging

Remain **available**  
despite failures



Meet stringent  
**performance** requirements



Keep data  
strongly **consistent**



# Failures and Availability

- Tens of thousands of servers
- Many types of failures
  - Power failures
  - Disk failures
  - OS crashes
  - Network failures
  - ...
- At data center scale, failures happen all the time!
- Services need to remain available!



# Facebook crippled by global outage across apps, employee systems



Downdetector said the Facebook outage is the largest it has seen, with more than 10.6 million reports worldwide. PHOTO: REUTERS

# Consistency

- User data is distributed across servers
  - Caching
  - Scalability
  - Tolerating failures
- Need to maintain data **consistency** and **persistence**
  - Data loss
  - Unexpected behaviors
- Very challenging in the presence of failures!

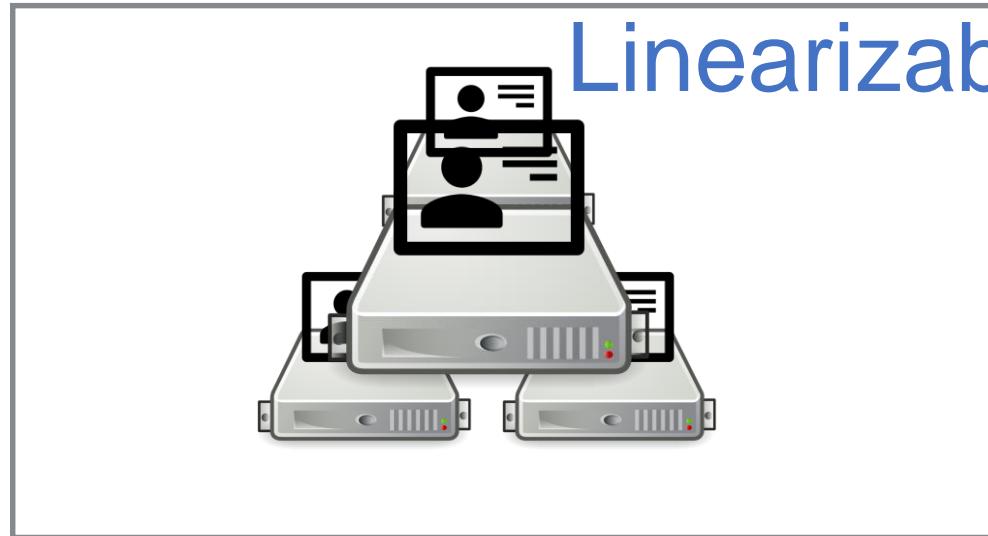


# Performance is Critical!

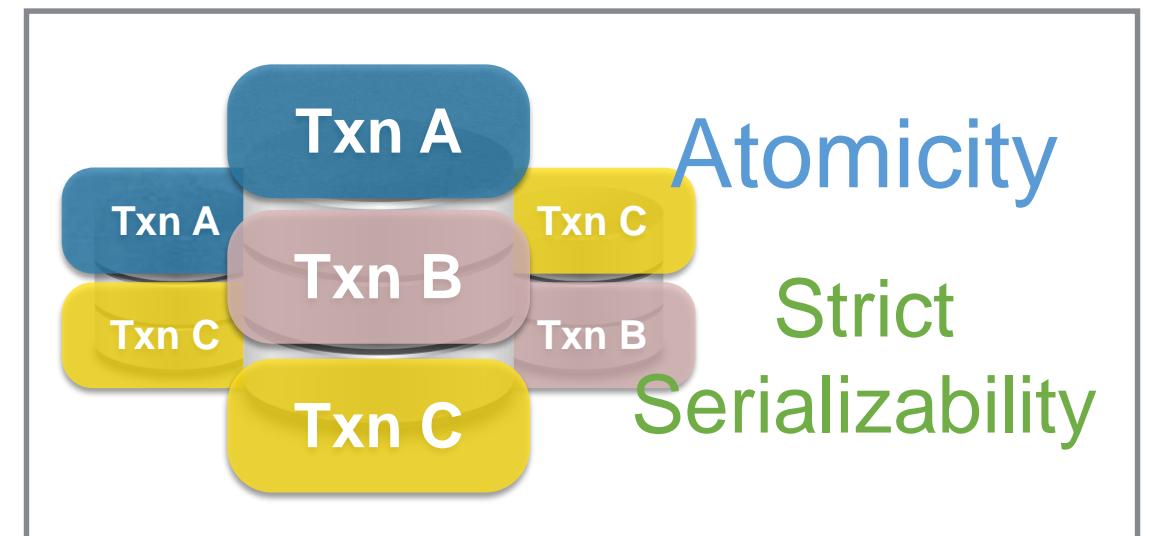
- $\uparrow$  service response time  
 $\downarrow$  revenue
- Each user request spawns into many sub-requests
- Response time dependent on the slowest subcomponent
- Tail latency matters!



# We Build System Software to Provide Fault Tolerance and Strong Consistency



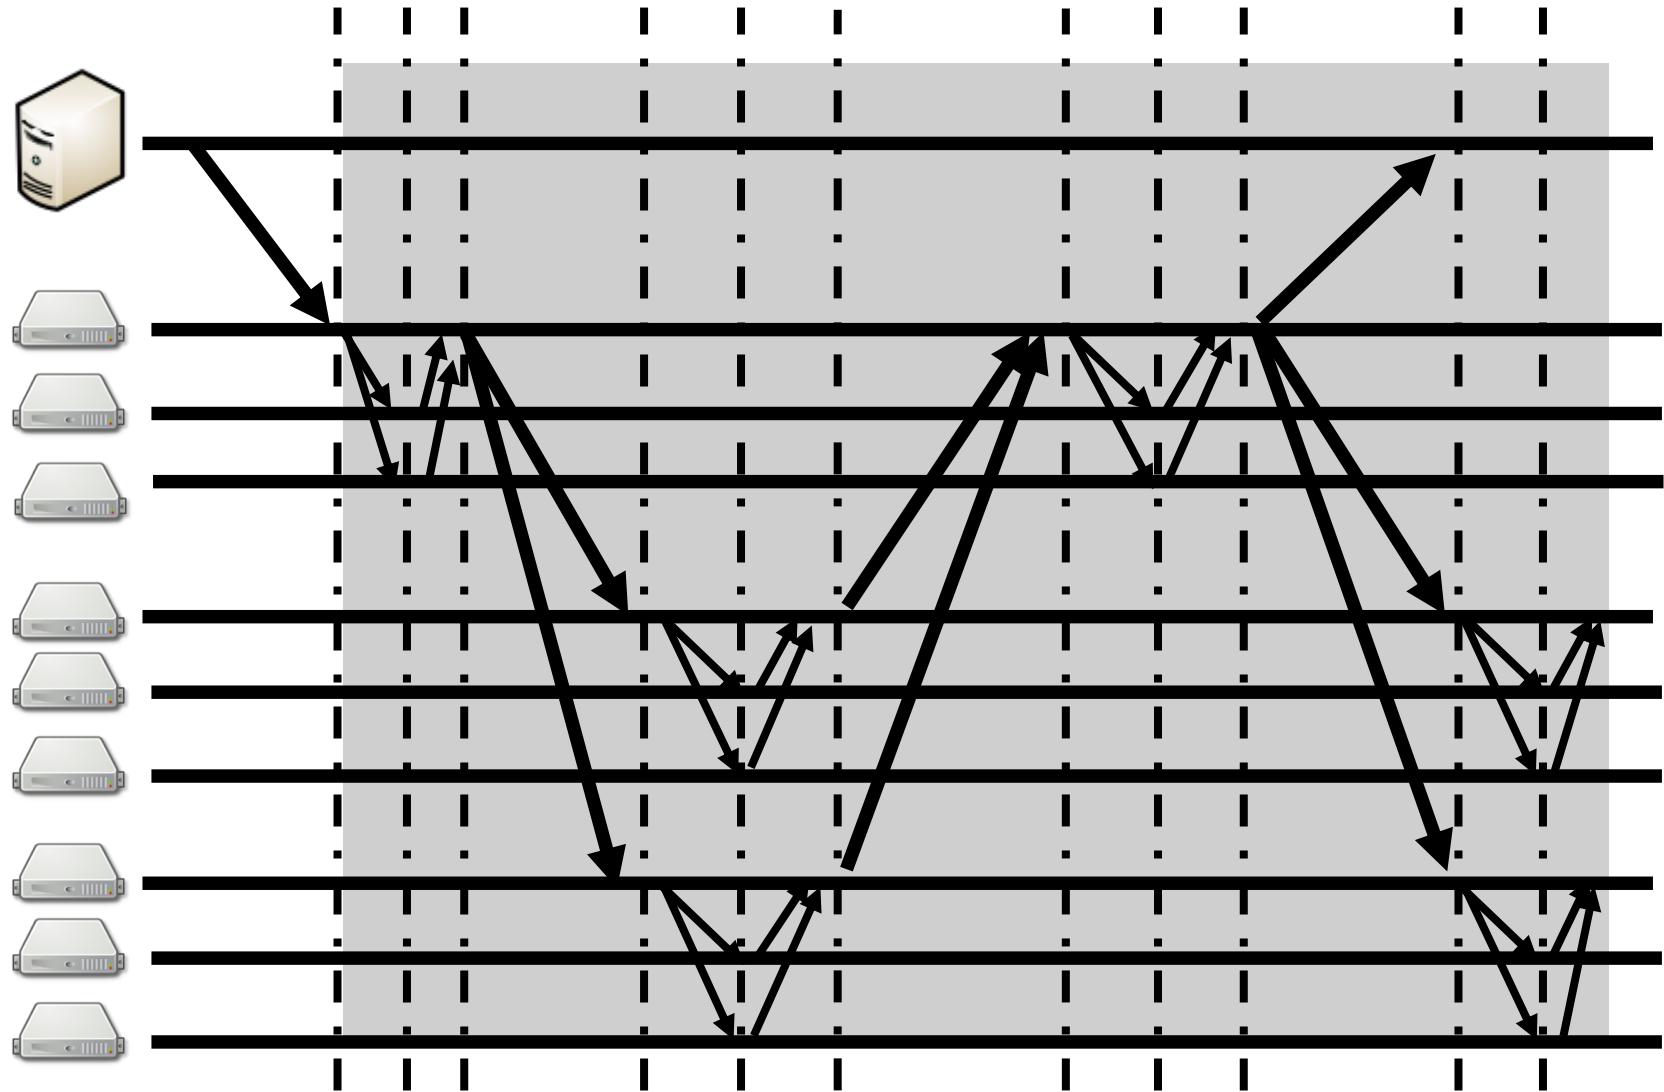
Fault-tolerant systems



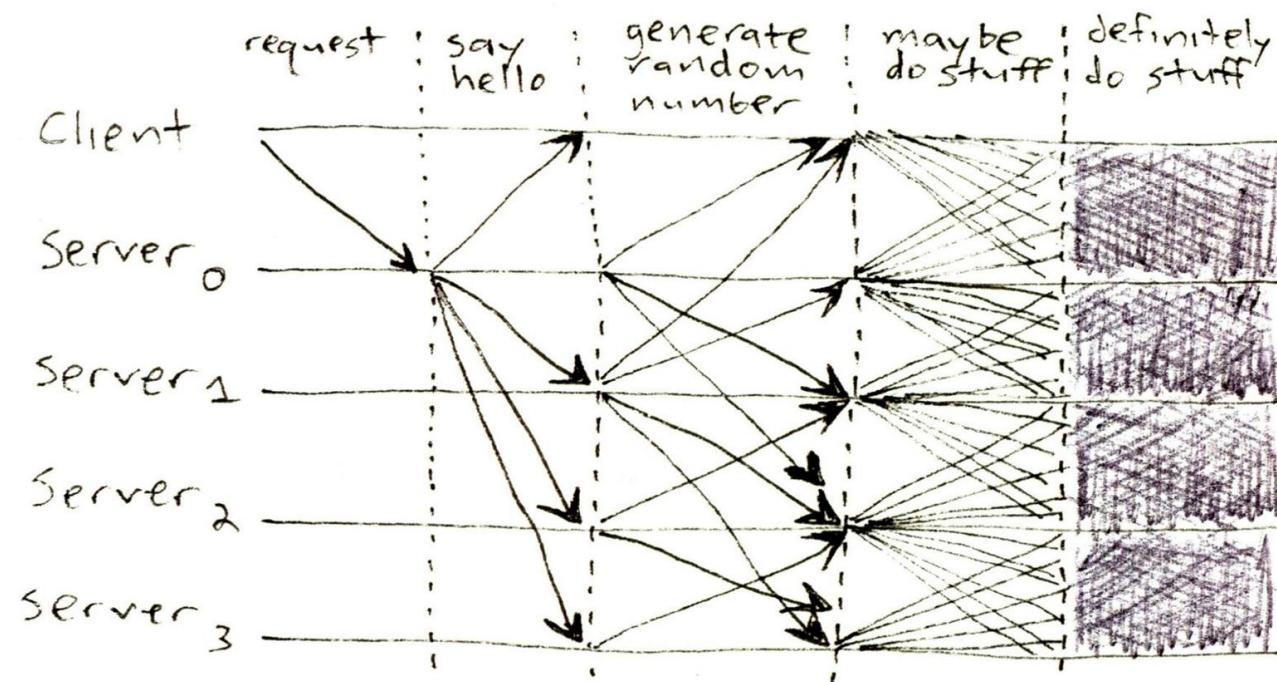
ACID transactions

# But... High Performance Cost!

Fault-Tolerant  
Distributed  
Transactions



# But... High Performance Cost!



[Mickens '13]

# Or... Weaker Guarantees for Better Performance



burden on application  
programmers



Amazon DynamoDB

We need **strong**  
guarantees!



Google Spanner

Can we build distributed systems with  
**both strong guarantees** and **high performance**?

Existing distributed systems designs  
consider the network

as a **black box**

# Current Network Model

Asynchrony: packets may arbitrarily

- Dropped
- Reordered
- Duplicated

receivers messages from A to B





# Trend: data center switches becoming more **programmable**

Programmable  
packet header

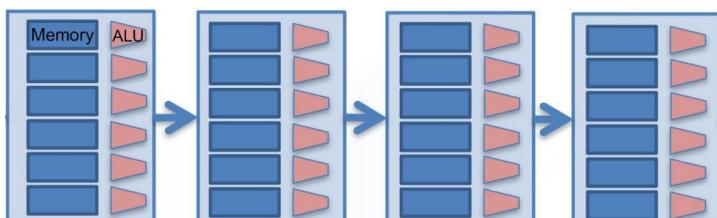
```
30 header_type ip;
31   fields {
32     version;
33     ihl : 4;
34     diffserv;
35     totalLength;
36     identifi;
37     flags;
38     fragOff;
39     ttl : 8;
40     protocol;
41     hdrCheck;
42     srcAddr : 32;
43     dstAddr : 32;
44   }
45 }
```

purpose  
memory

application-specific computation at **line rate**



Flexible packet  
processing pipeline

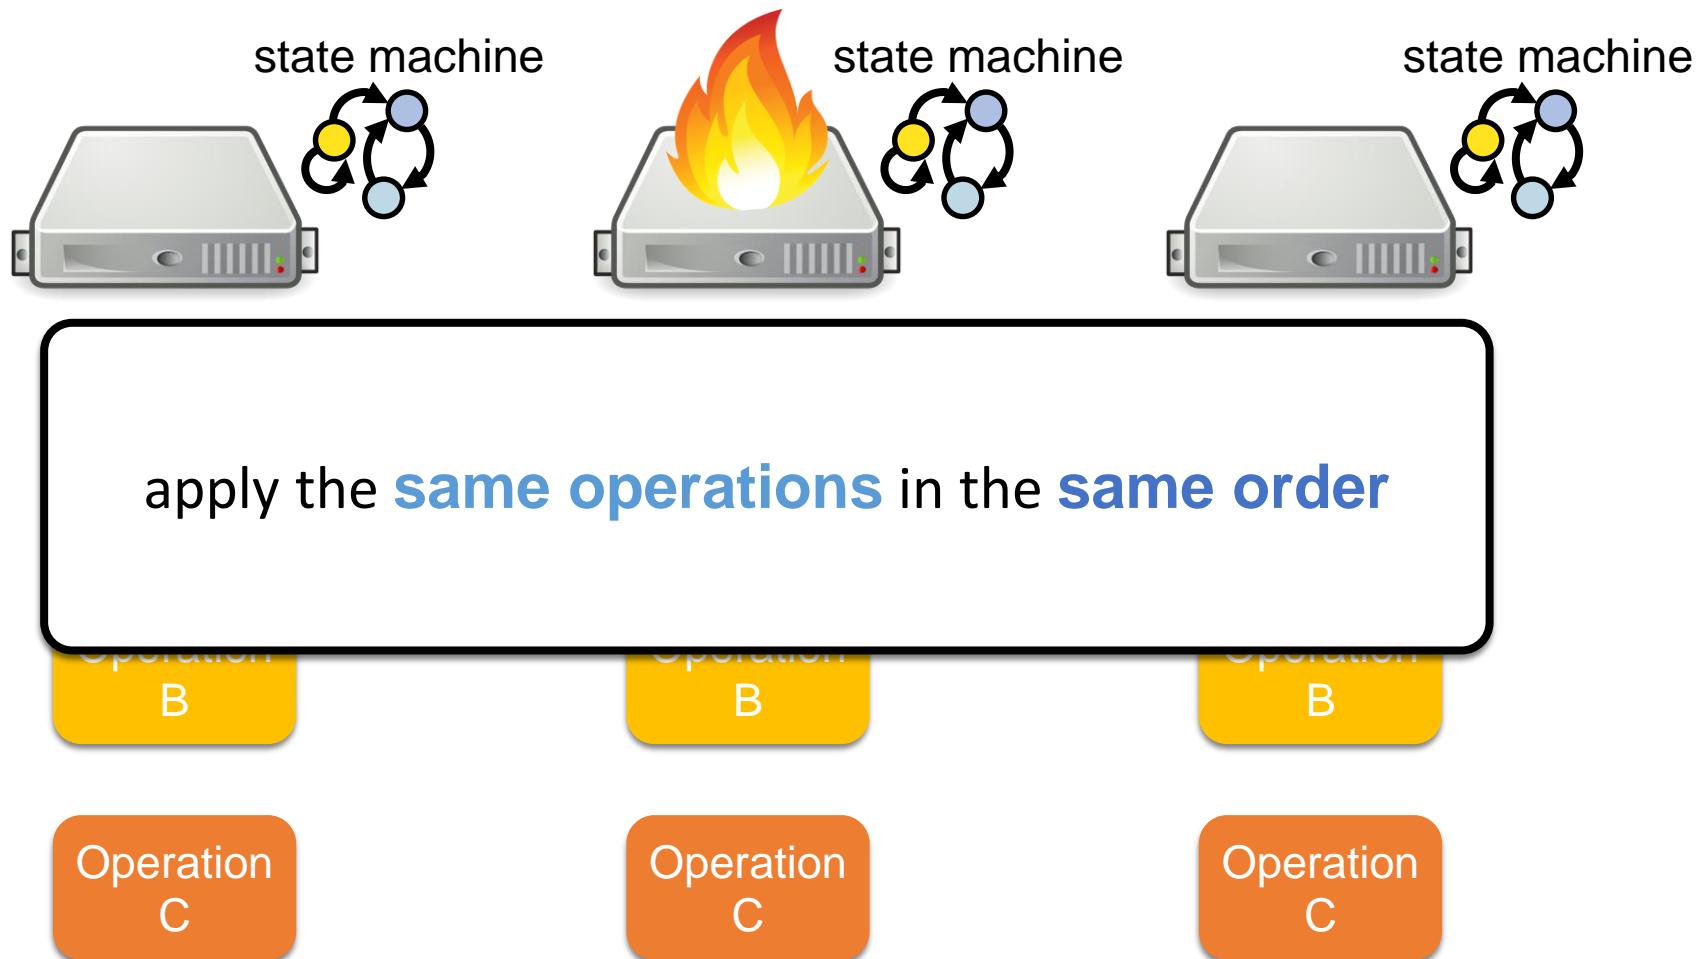


# New approach: co-designing distributed systems with the network

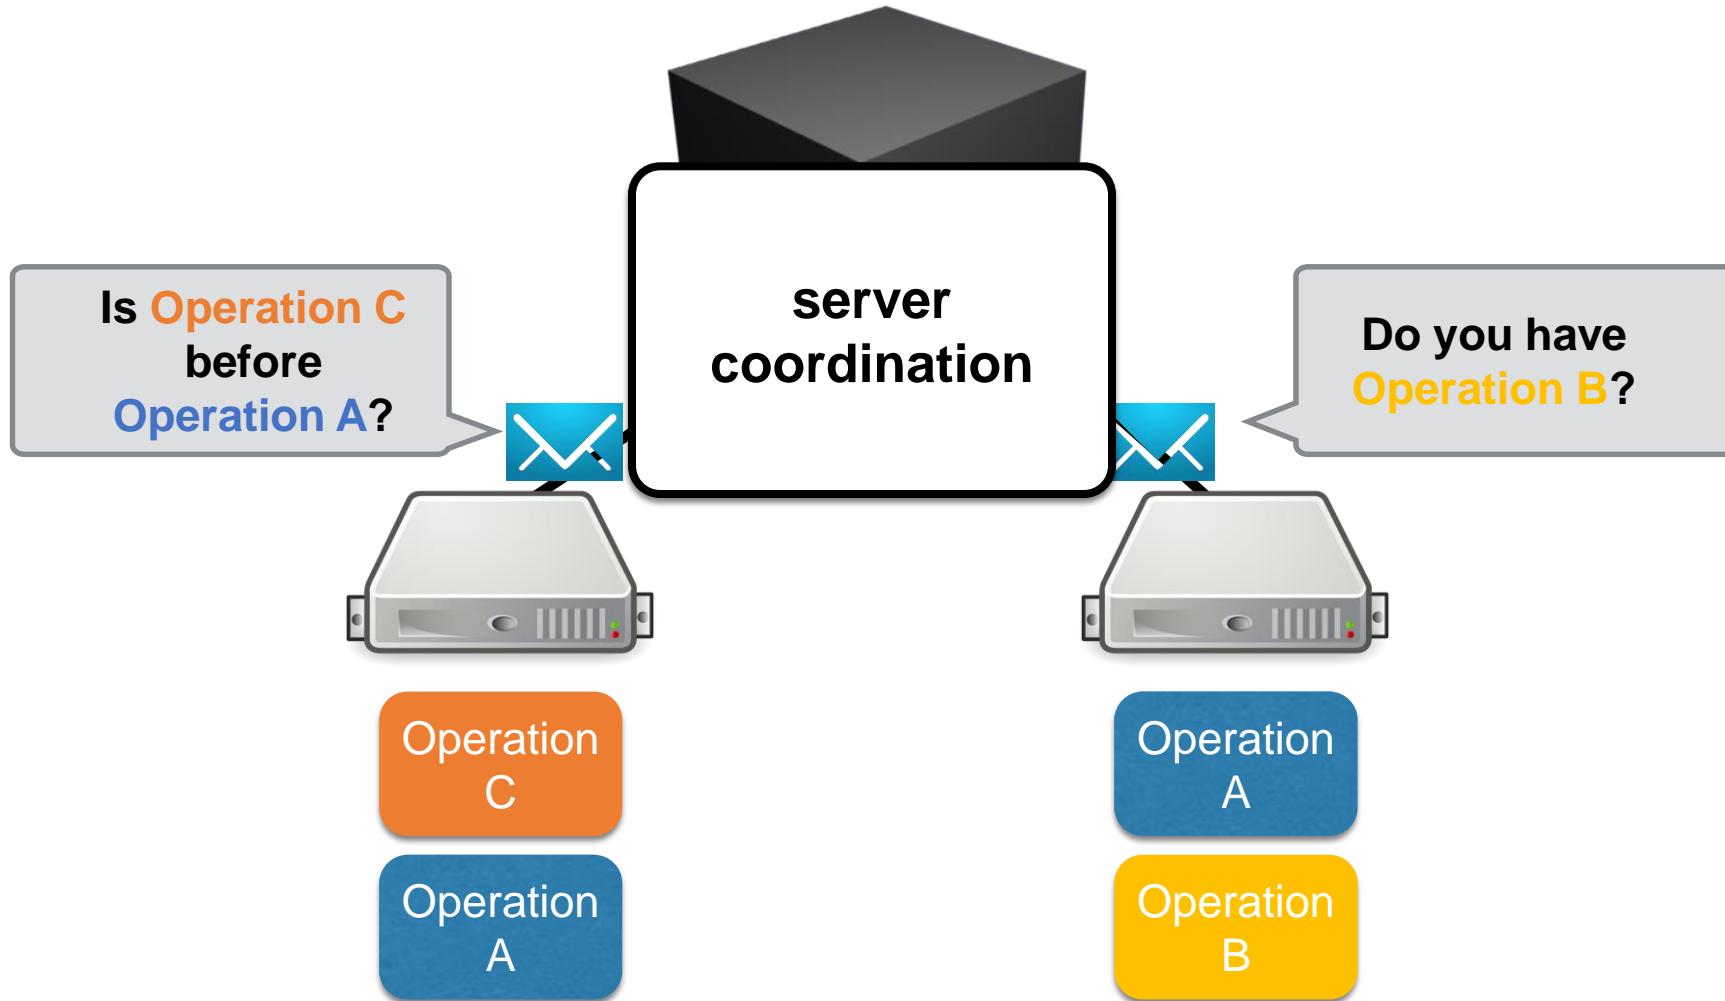
- No longer treat network as a black box
- Leverage programmable network devices to design stronger **network primitives**
- Build new **distributed protocols** around these network primitives
- Result: practical distributed systems with both **strong guarantees and high performance**
- No more weak consistencies!

# Co-Design Approach Applied to State Machine Replication

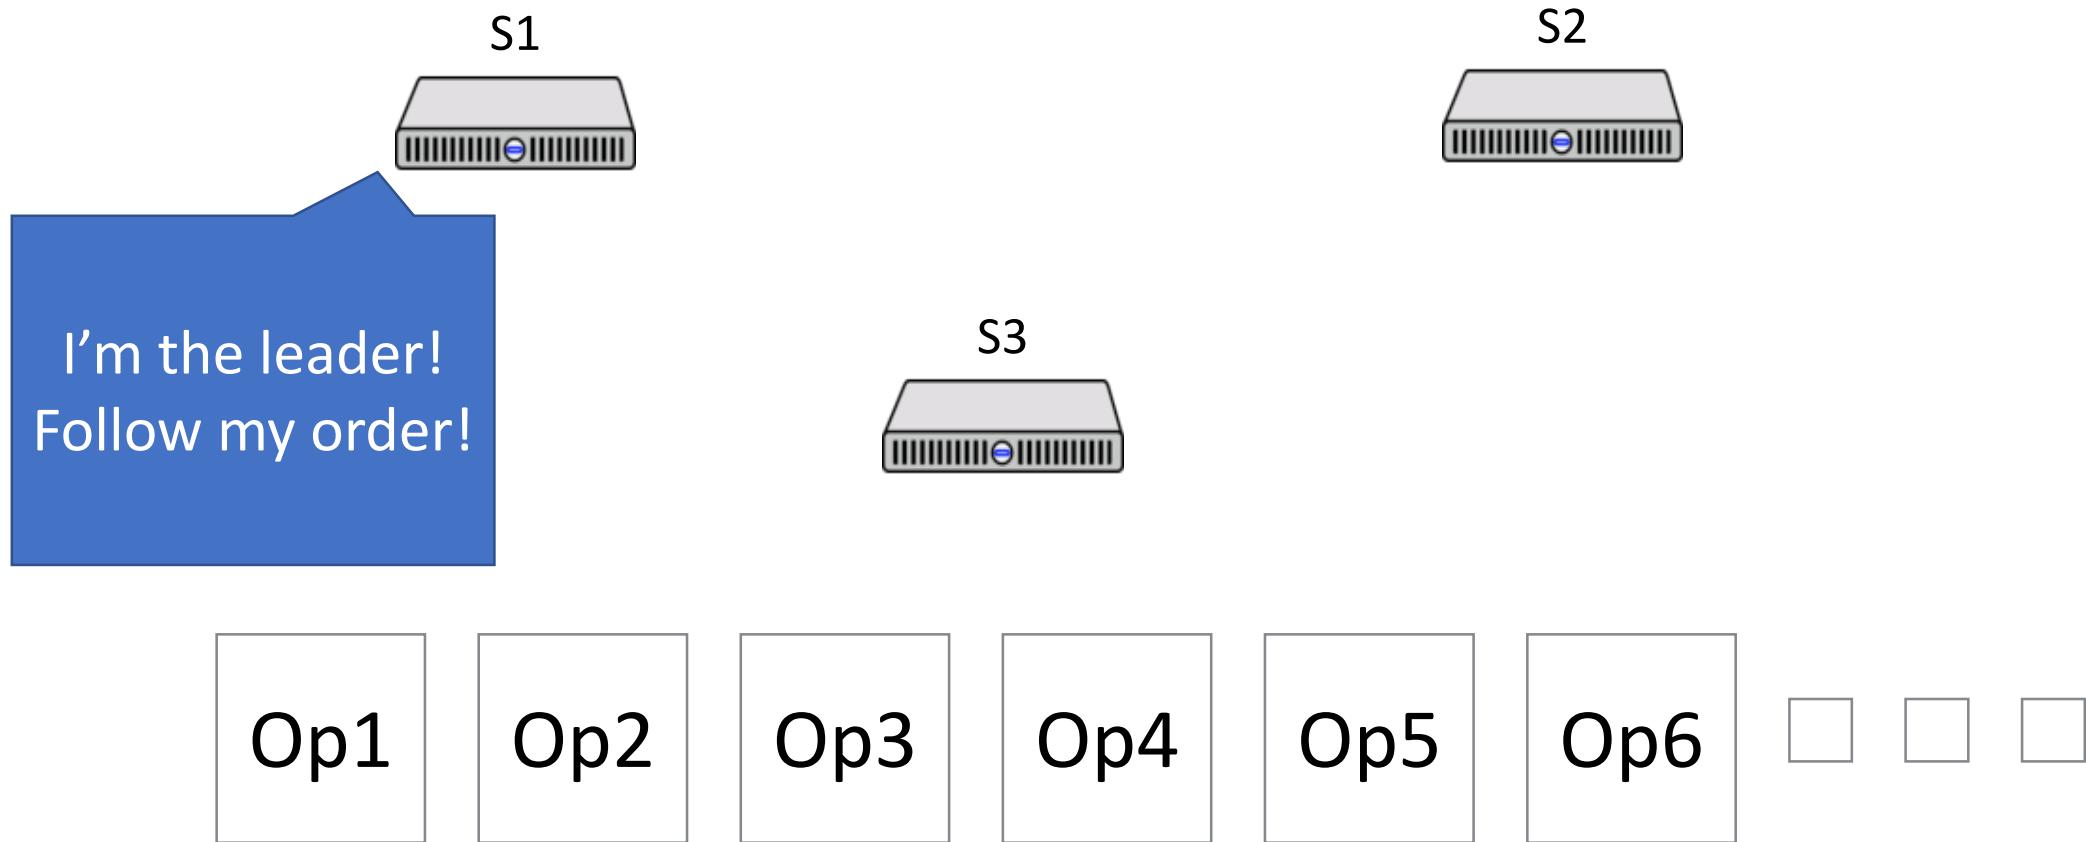
# State Machine Replication Primer



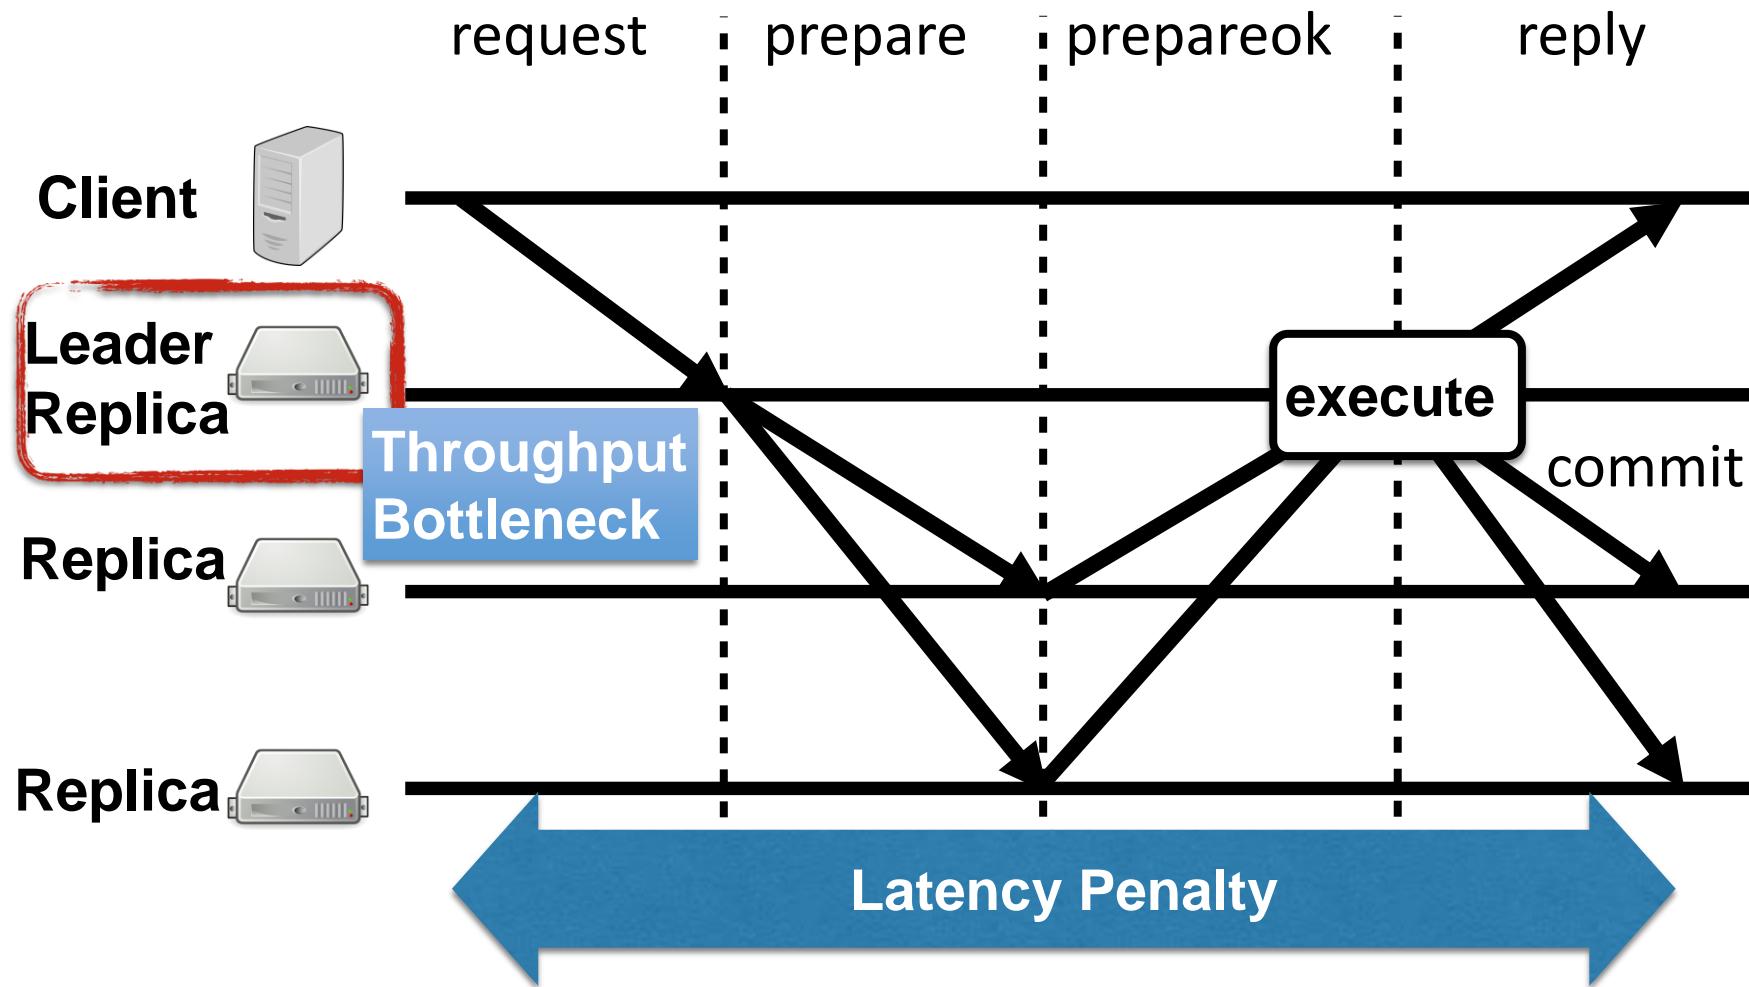
# Black-box network model requires server coordination



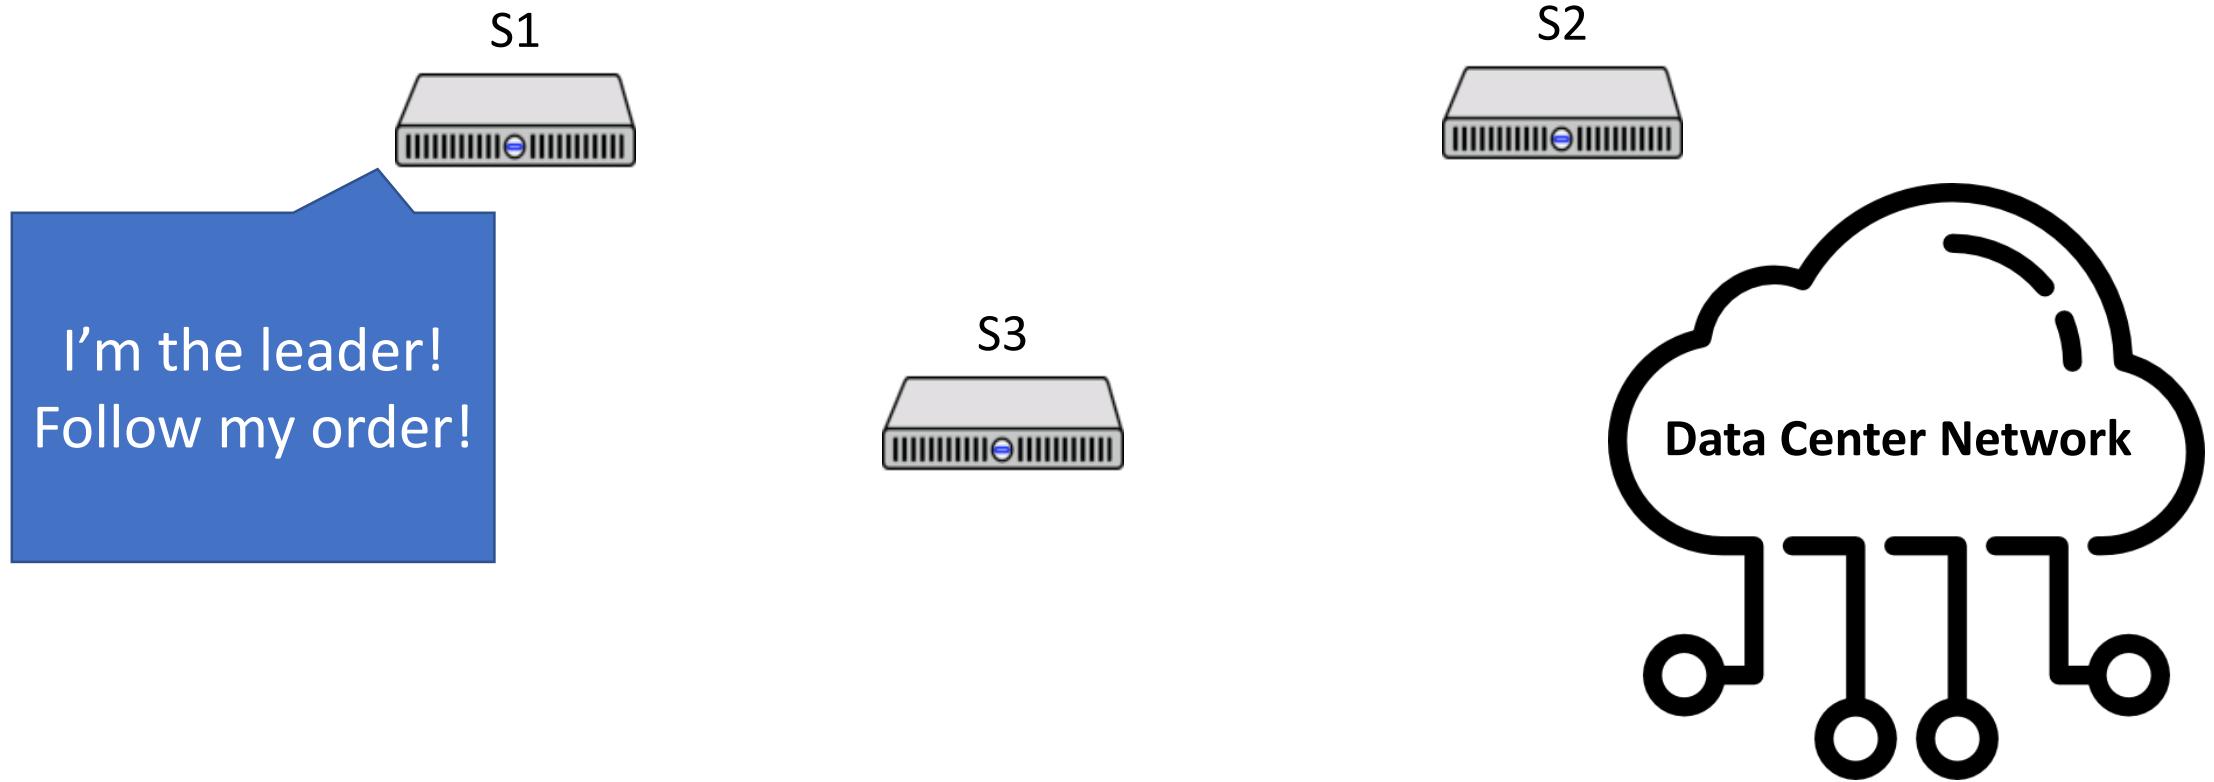
# How do we ensure same order of requests?



# Standard approach: Paxos/VR/Raft



# What if... the network orders requests for us?



# New network primitive: Ordered Unreliable Multicast

- **Multicast:** client sends a single message. network delivers the message to all replicas
- **Ordered multicast:** all replicas receive multicast messages **in the same order**
- **Unreliable delivery:** no guarantee that messages will be delivered
- **Drop detection:** replicas receive **notifications** for dropped messages

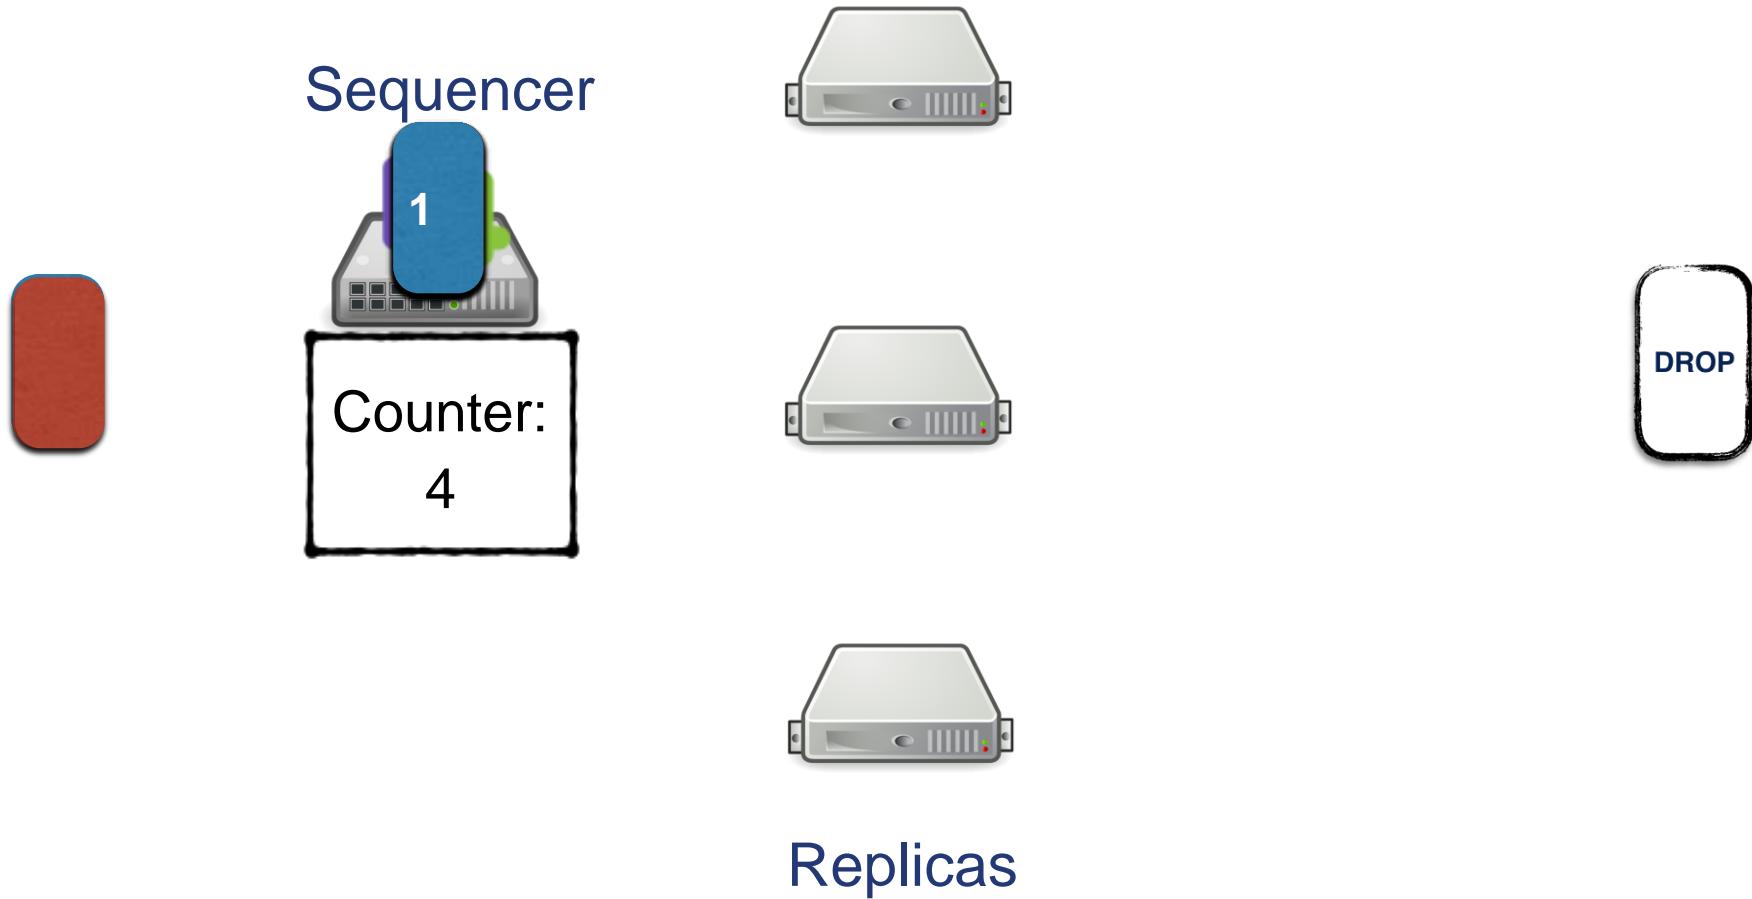
# How to implement ordering in the network??



Client requests first go through  
a programmable switch

Switch writes a **counter** value  
into request packet

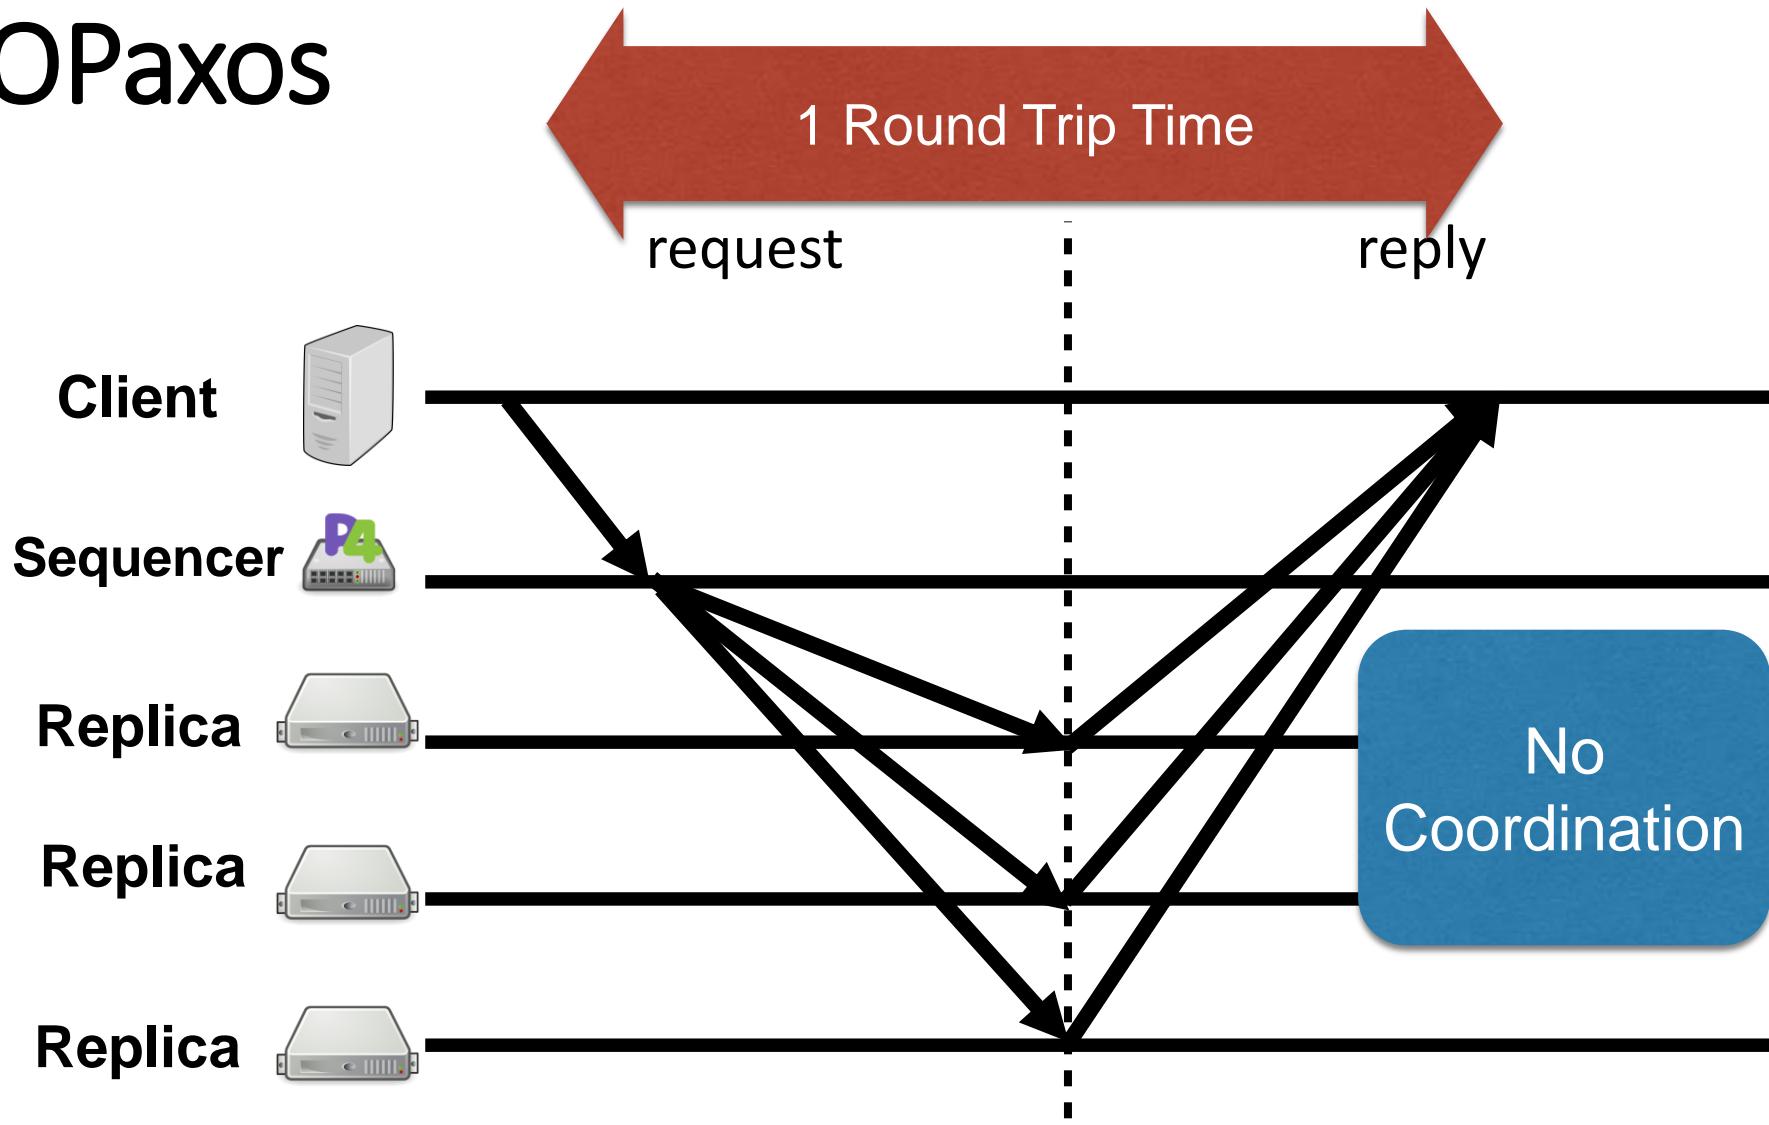
# How to implement Ordered Unreliable Multicast (efficiently)?



# How does Ordered Unreliable Multicast benefit state machine replication?

- All replicas receive messages in the **same order**
- If they receive all the messages
- Just execute and reply
  - they **know** other replicas also receive the messages in order
- All replicas are in **consistent state**
- **No coordination** at all!

# Co-designing SMR with Network Ordering: NOPaxos

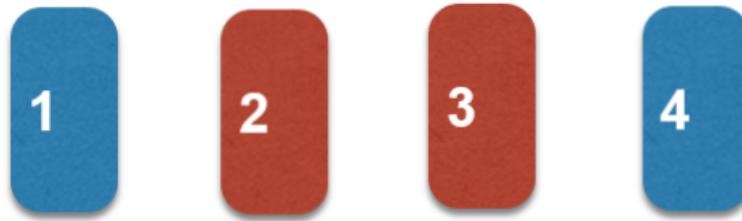
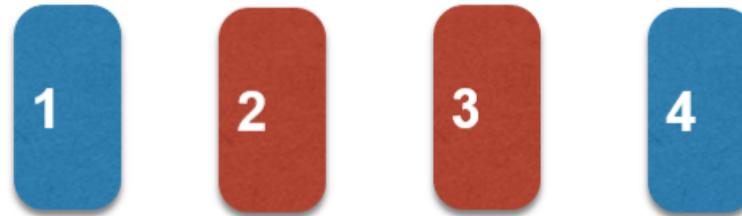


What happens when messages  
are dropped in the network?

Sequencer



Replicas



# Gap Agreement

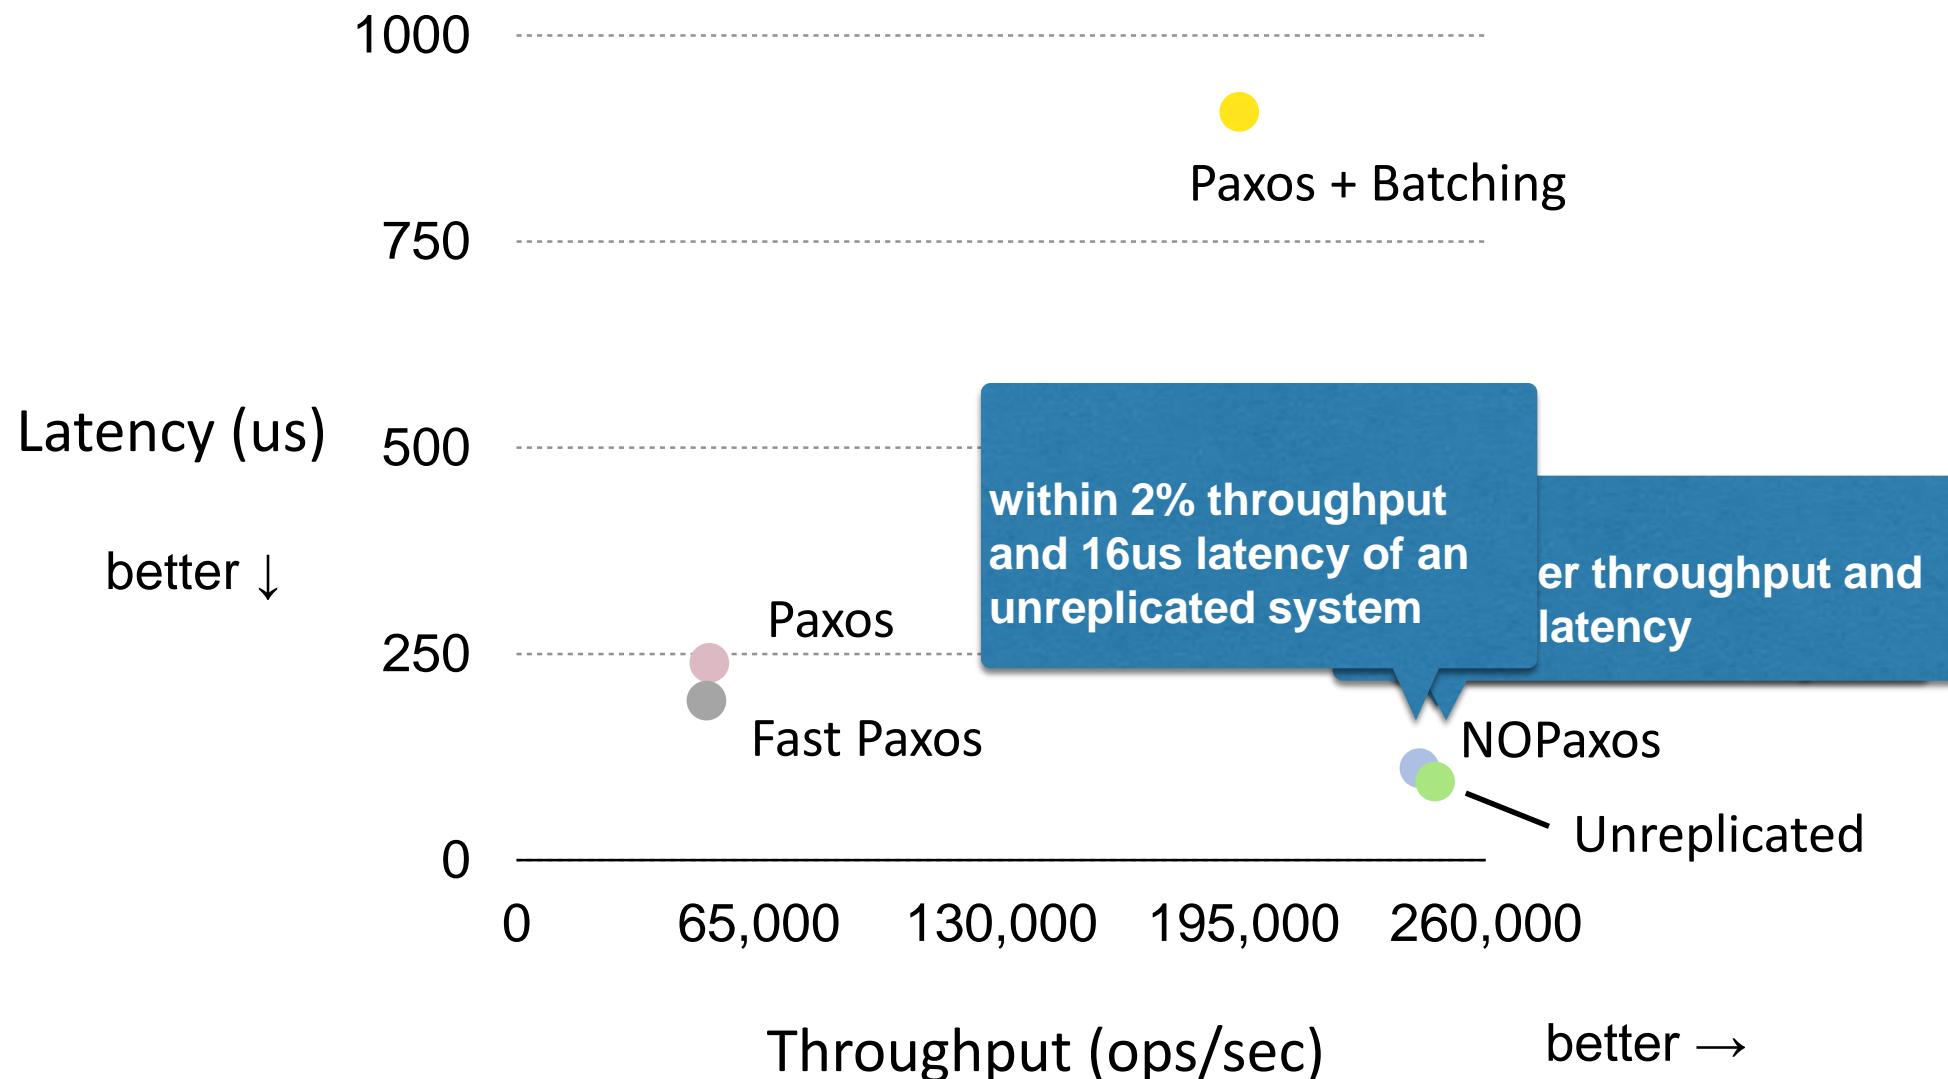
- Need to agree if the message is **received or permanently dropped**
- When a replica detects message drops
  - **Non-leader replicas:** recover the missing message from the leader
  - **Leader replica:** coordinates to commit a NO-OP (Paxos)
- Efficient recovery from network anomalies

# Other protocols

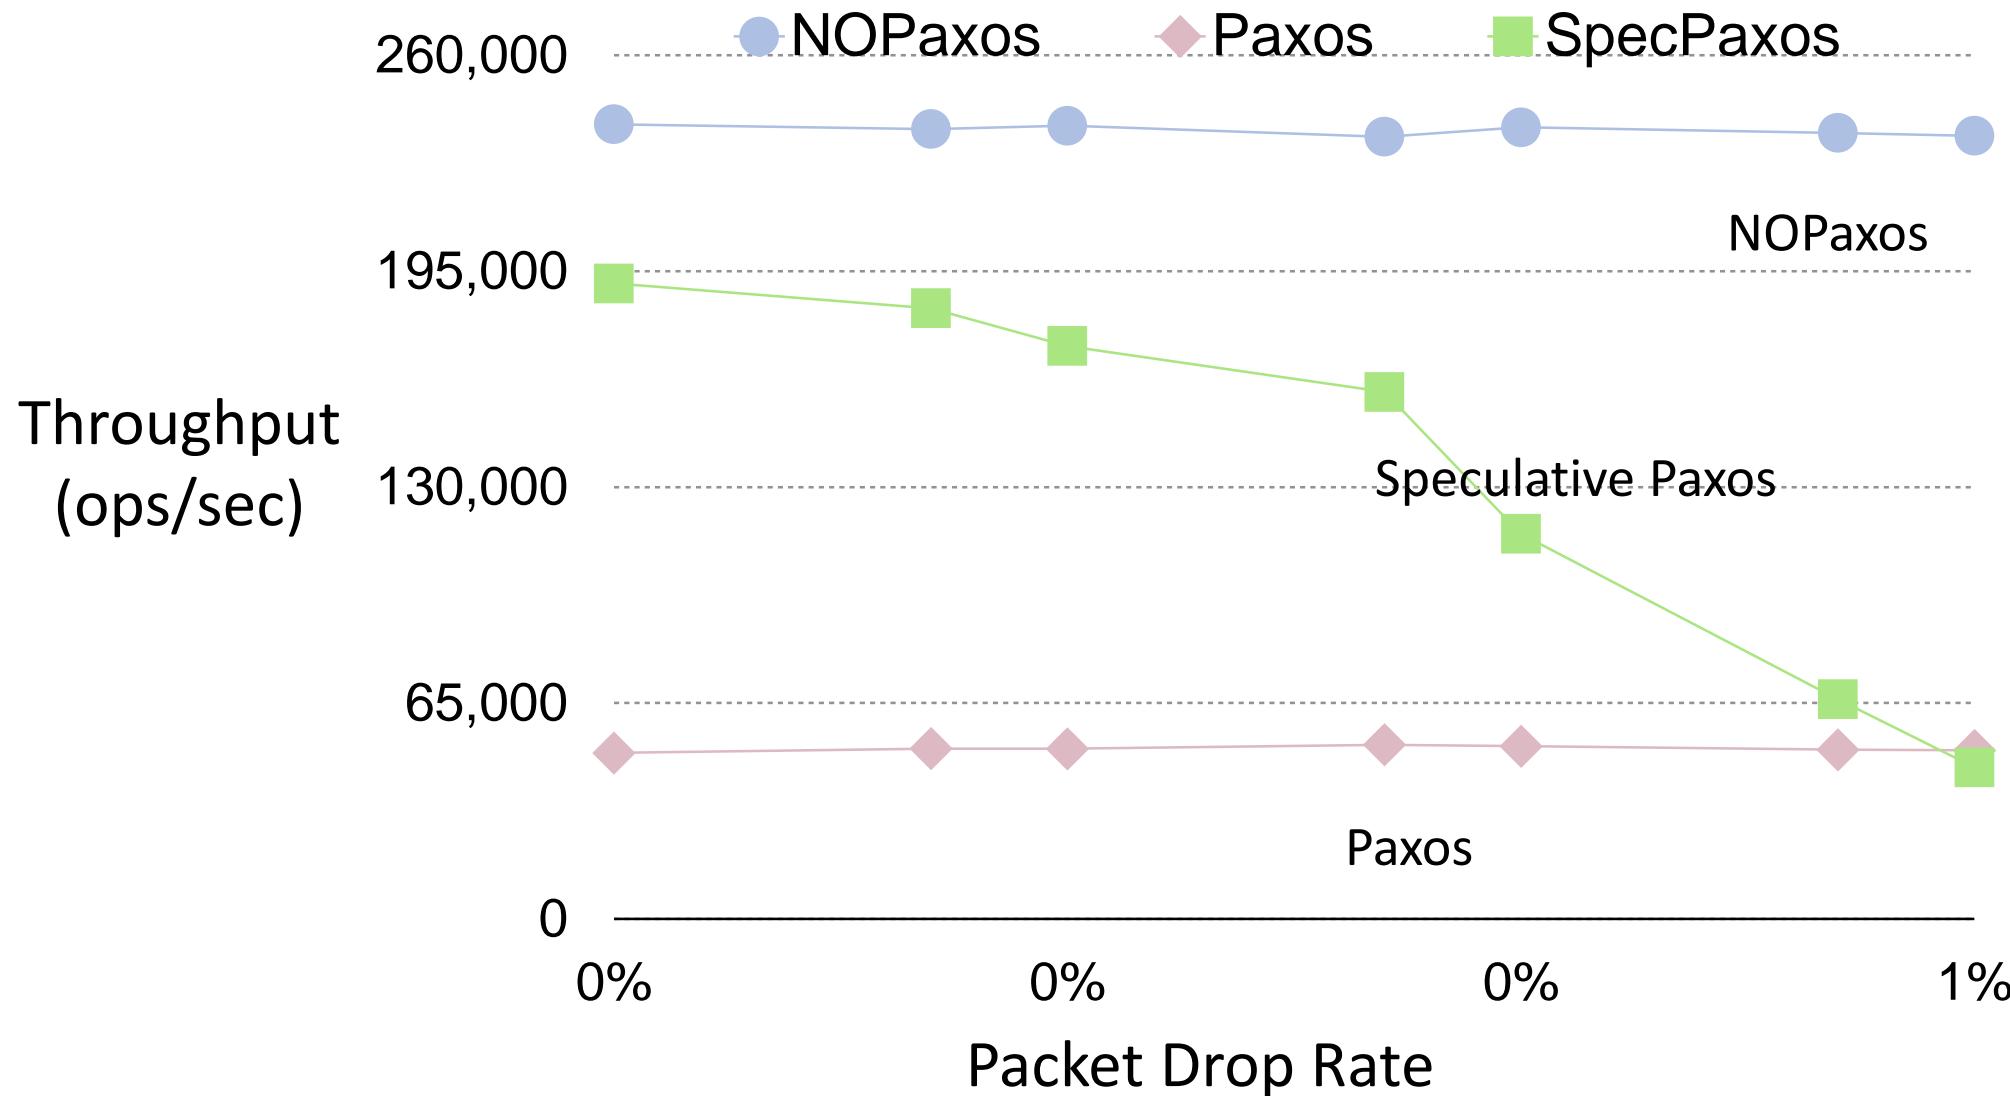
- Handling leader replica failures
- Handling sequencer switch failures
- More details in the paper!

What benefits do we get?

# NOPaxos has no Paxos overhead!



# NOPaxos is resilient to network anomalies

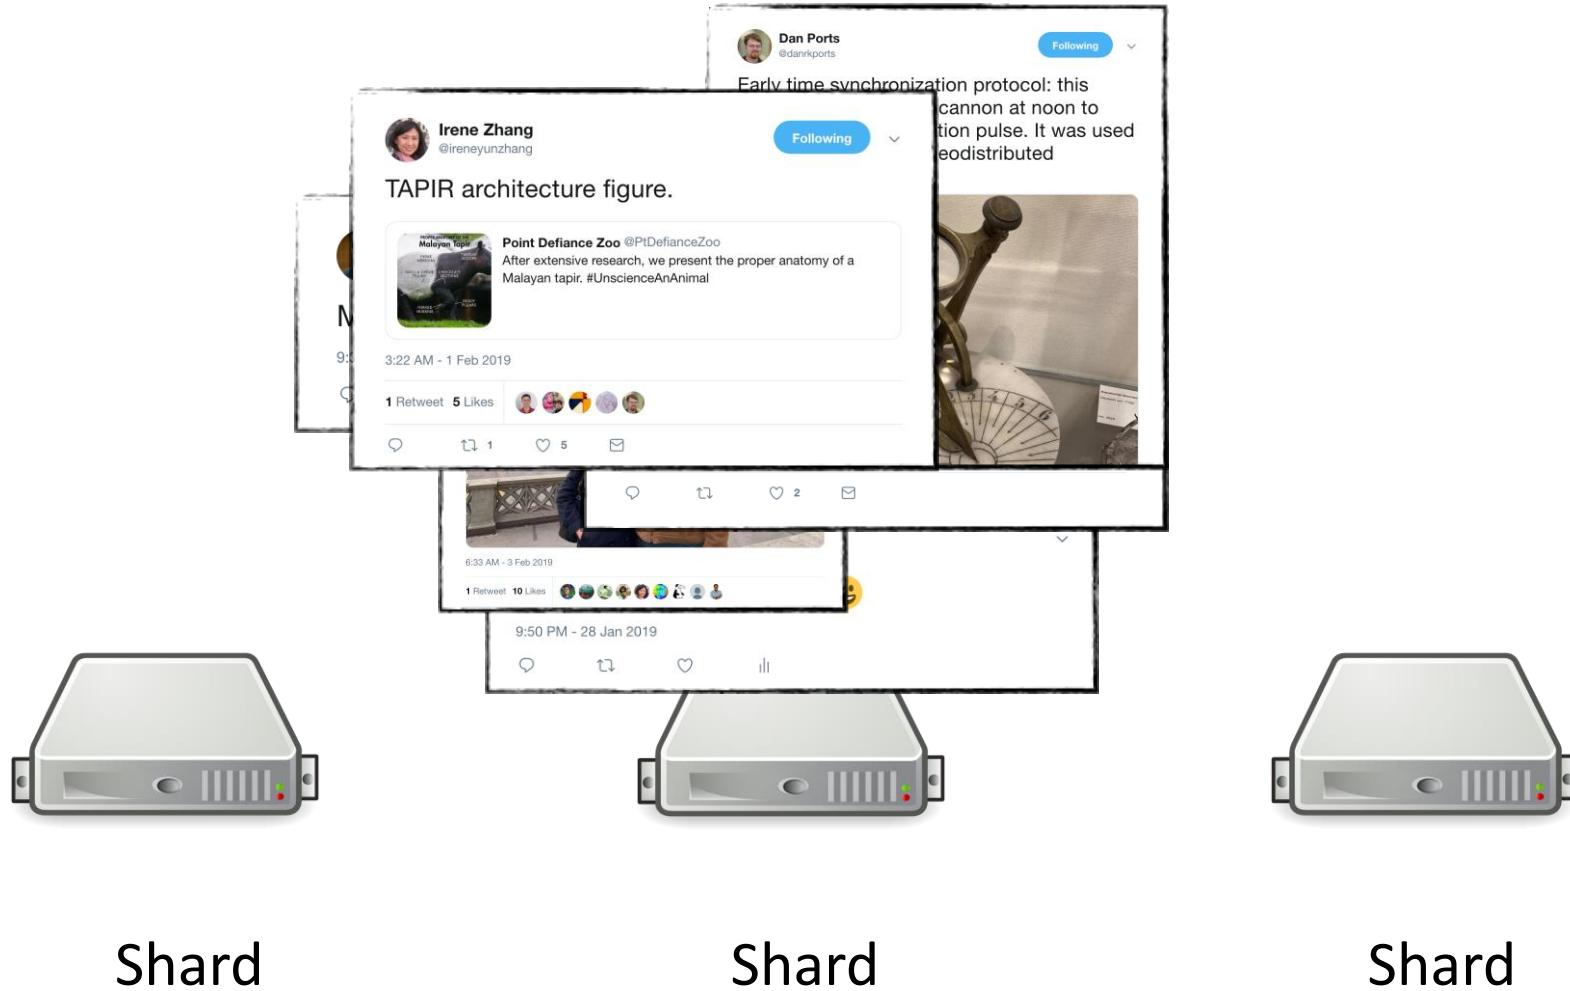


What else can do with this  
Co-Design approach?

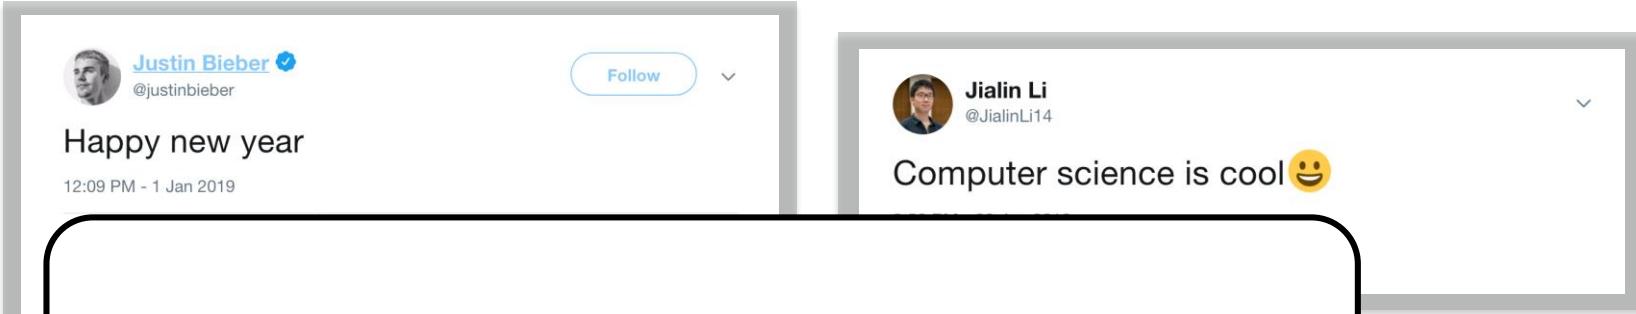
Many real-workloads are  
**skewed** and **dynamic**



# Data partitioning for scalability



# Skewed workloads lead to load imbalance



popular objects can **overload** storage servers



# Skewed workloads are **diverse**

- read-heavy
- write-heavy
- read-write mixed

The screenshot shows a dark purple header with the OSDI'20 logo and navigation links: ATTEND, PROGRAM, PARTICIPATE, SPONSORS, and ABOUT. To the right is a blue Twitter icon. Below the header is a white slide area containing the text: "A large scale analysis of hundreds of in-memory cache clusters at Twitter".



**Workload Analysis of a Large-Scale Key-Value Store**

- small objects
- large objects
- combination of both



**Fast key-value stores: An idea whose time has come and gone**

# Two approaches to deal with highly skewed workloads

## Caching

- Cache popular objects in a faster tier
- Caching tier *absorbs* traffic to popular objects
- More uniform load on backend storage servers



Caching Layer



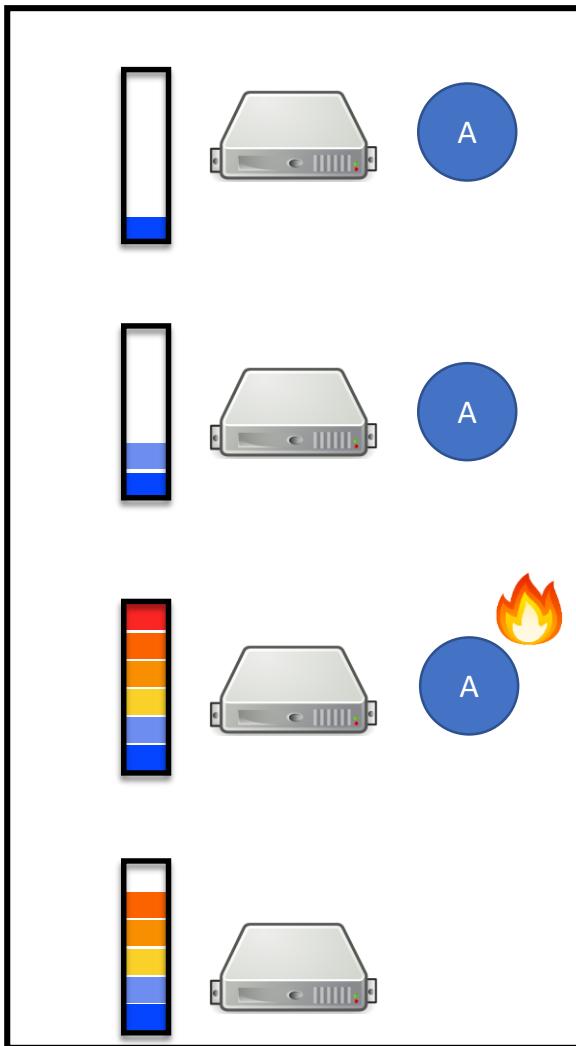
# Two approaches to deal with highly skewed workloads



## Selective Replication

- **Replicate popular objects on multiple servers**
- Requests to replicated objects can be forwarded to *any* replica
- **Distribute** load across servers

# Selective Replication is challenging



How to route requests to  
the right server?

How to ensure  
consistency?

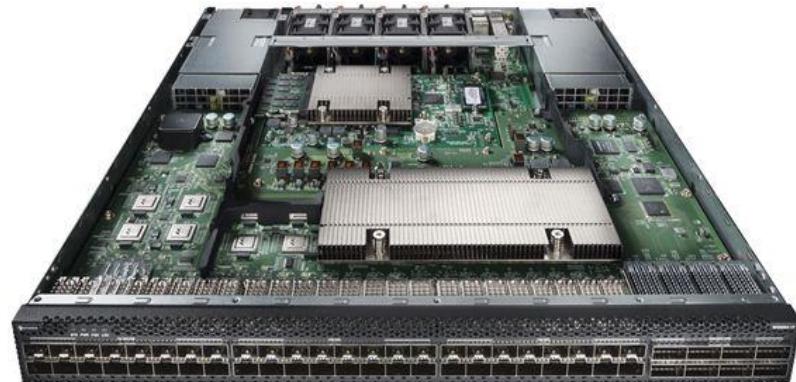
What if the network can track  
where the copies are?

# Pegasus' approach

selective  
replication

+

in-network  
**coherence directory**

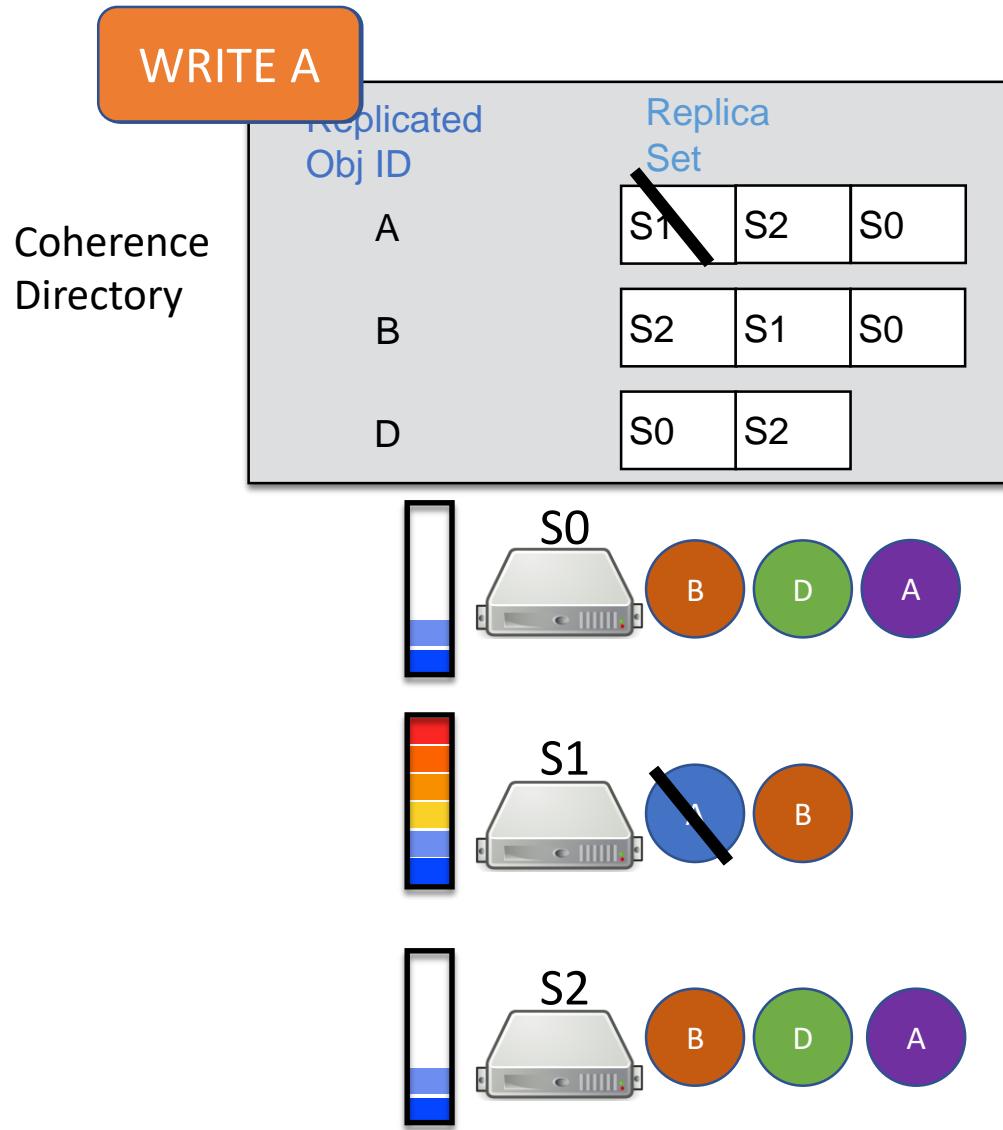


Barefoot P4 switch

# Coherence directory for replicated data

- Inspired by CPU cache coherence protocols
- Centralized directory that **tracks**:
  - **Which** objects are replicated
  - **Location** of replicated objects
- **Forwards** requests to server with spare capacity
- Ensures **strong consistency**

# Coherence directory illustrated



# Implementing coherence directory in the network

rack-scale storage system

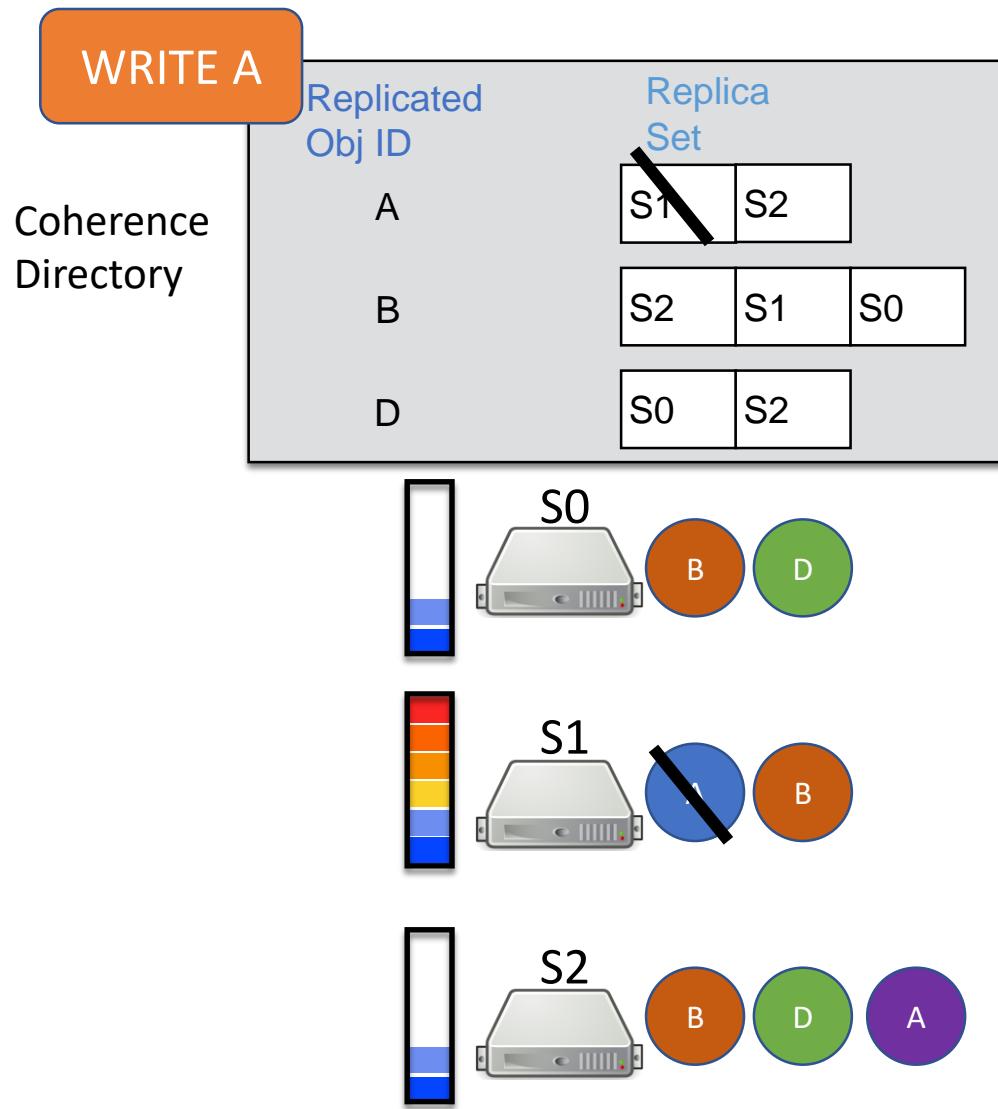
- All requests and replies traverse the ToR switch
- ToR serves as a **central point**
- **Line-rate** packet processing
  - No throughput bottleneck
  - Zero latency overhead



# How to ensure consistency?

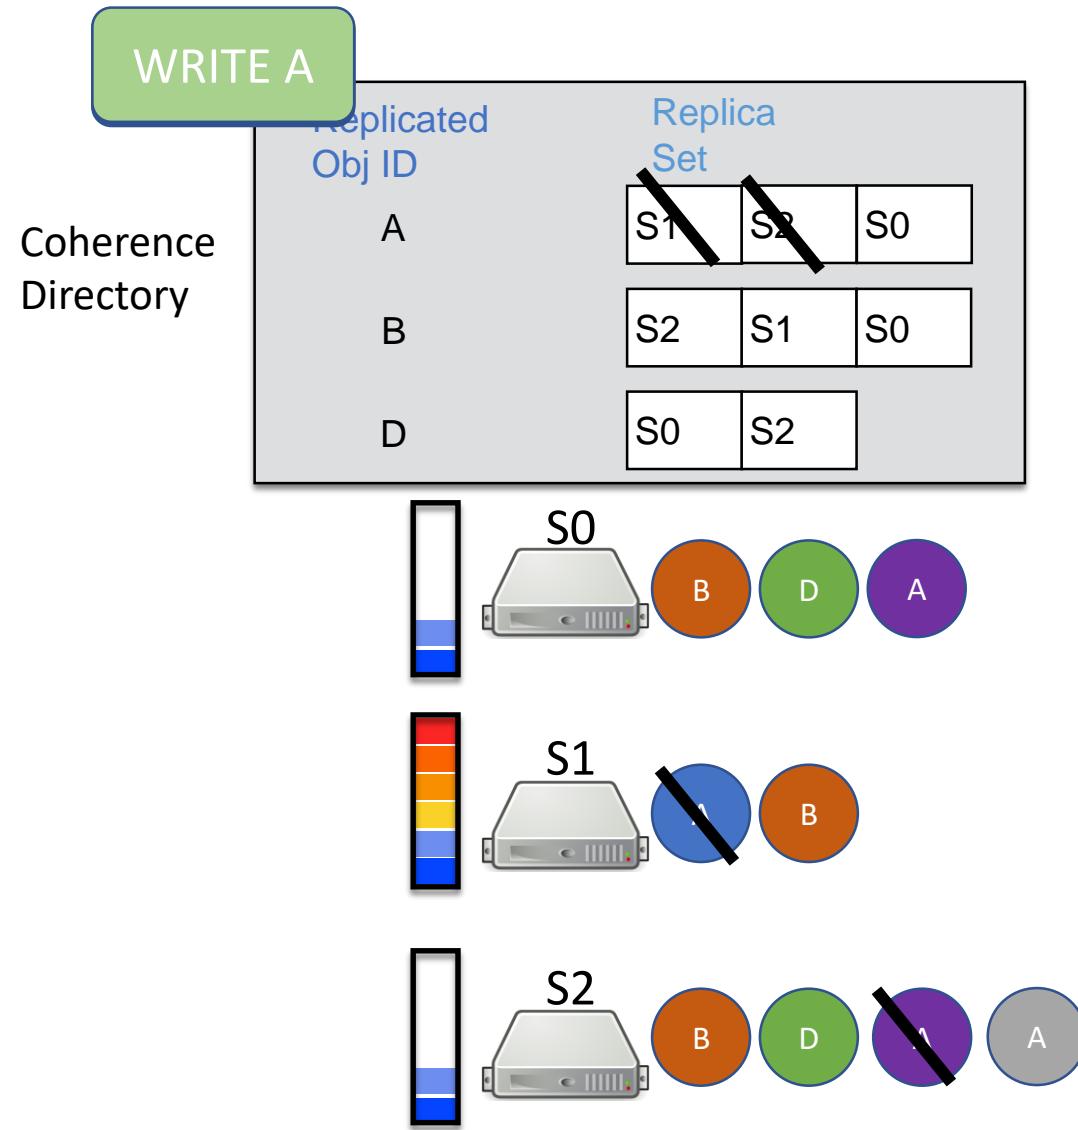
- Switch processes **all** requests
- Switch directory tracks which servers have the **latest** copy
- When do you update the directory?
  - Receive request?
  - Receive reply?

# When do you update the directory?



How to deal with **asynchrony** in  
the system?

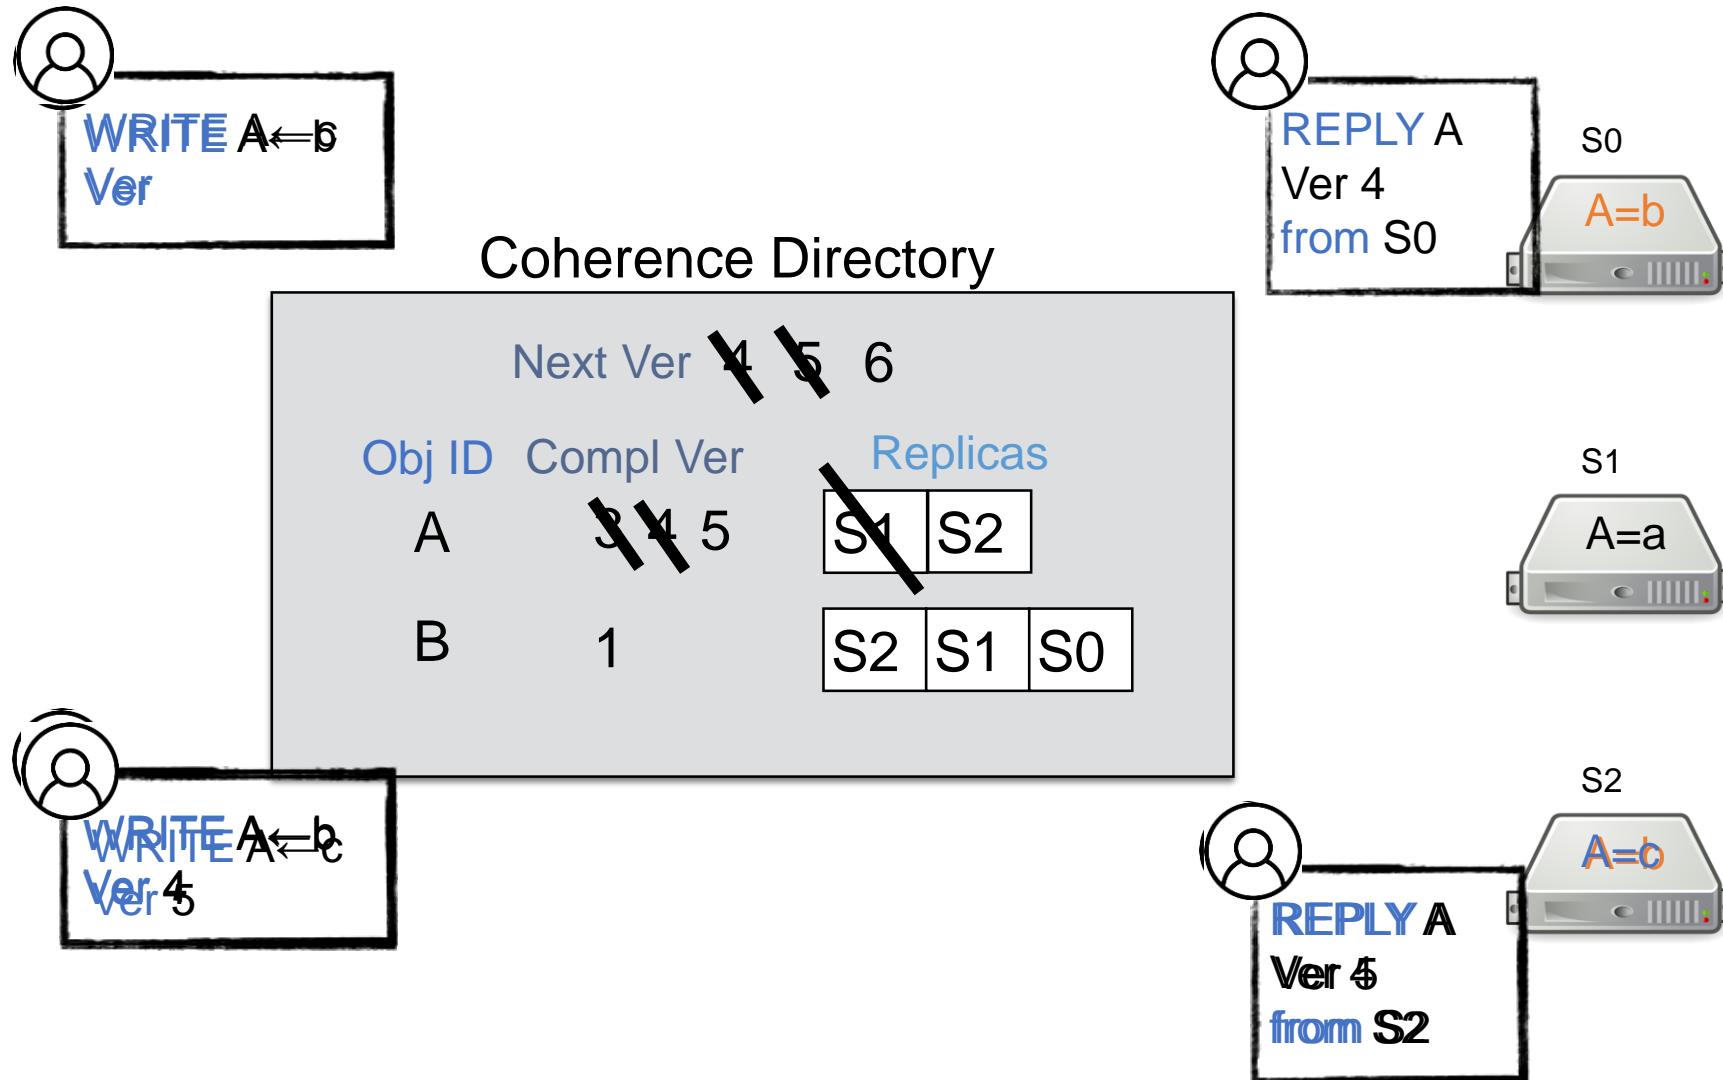
# Issues with Asynchrony



# How to deal with asynchrony?

Version-based  
coherence protocol!

# Pegasus' coherence protocol in action

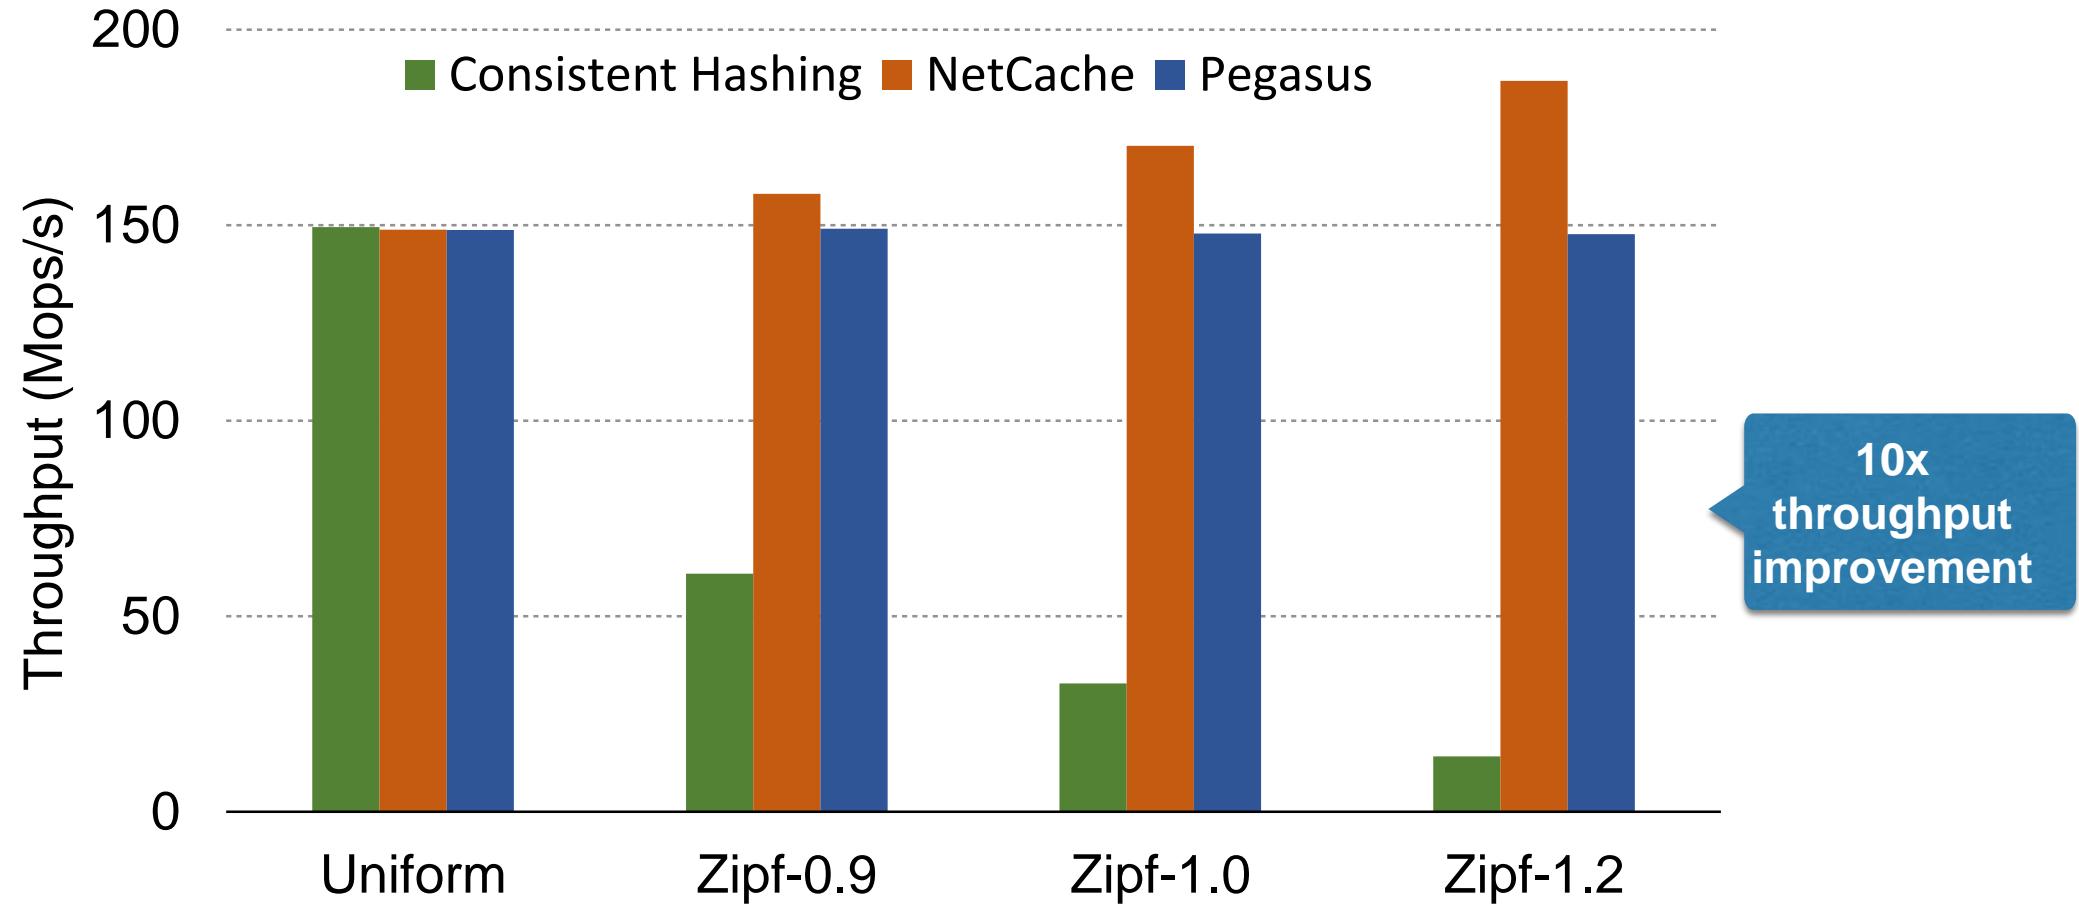


# Other protocol details

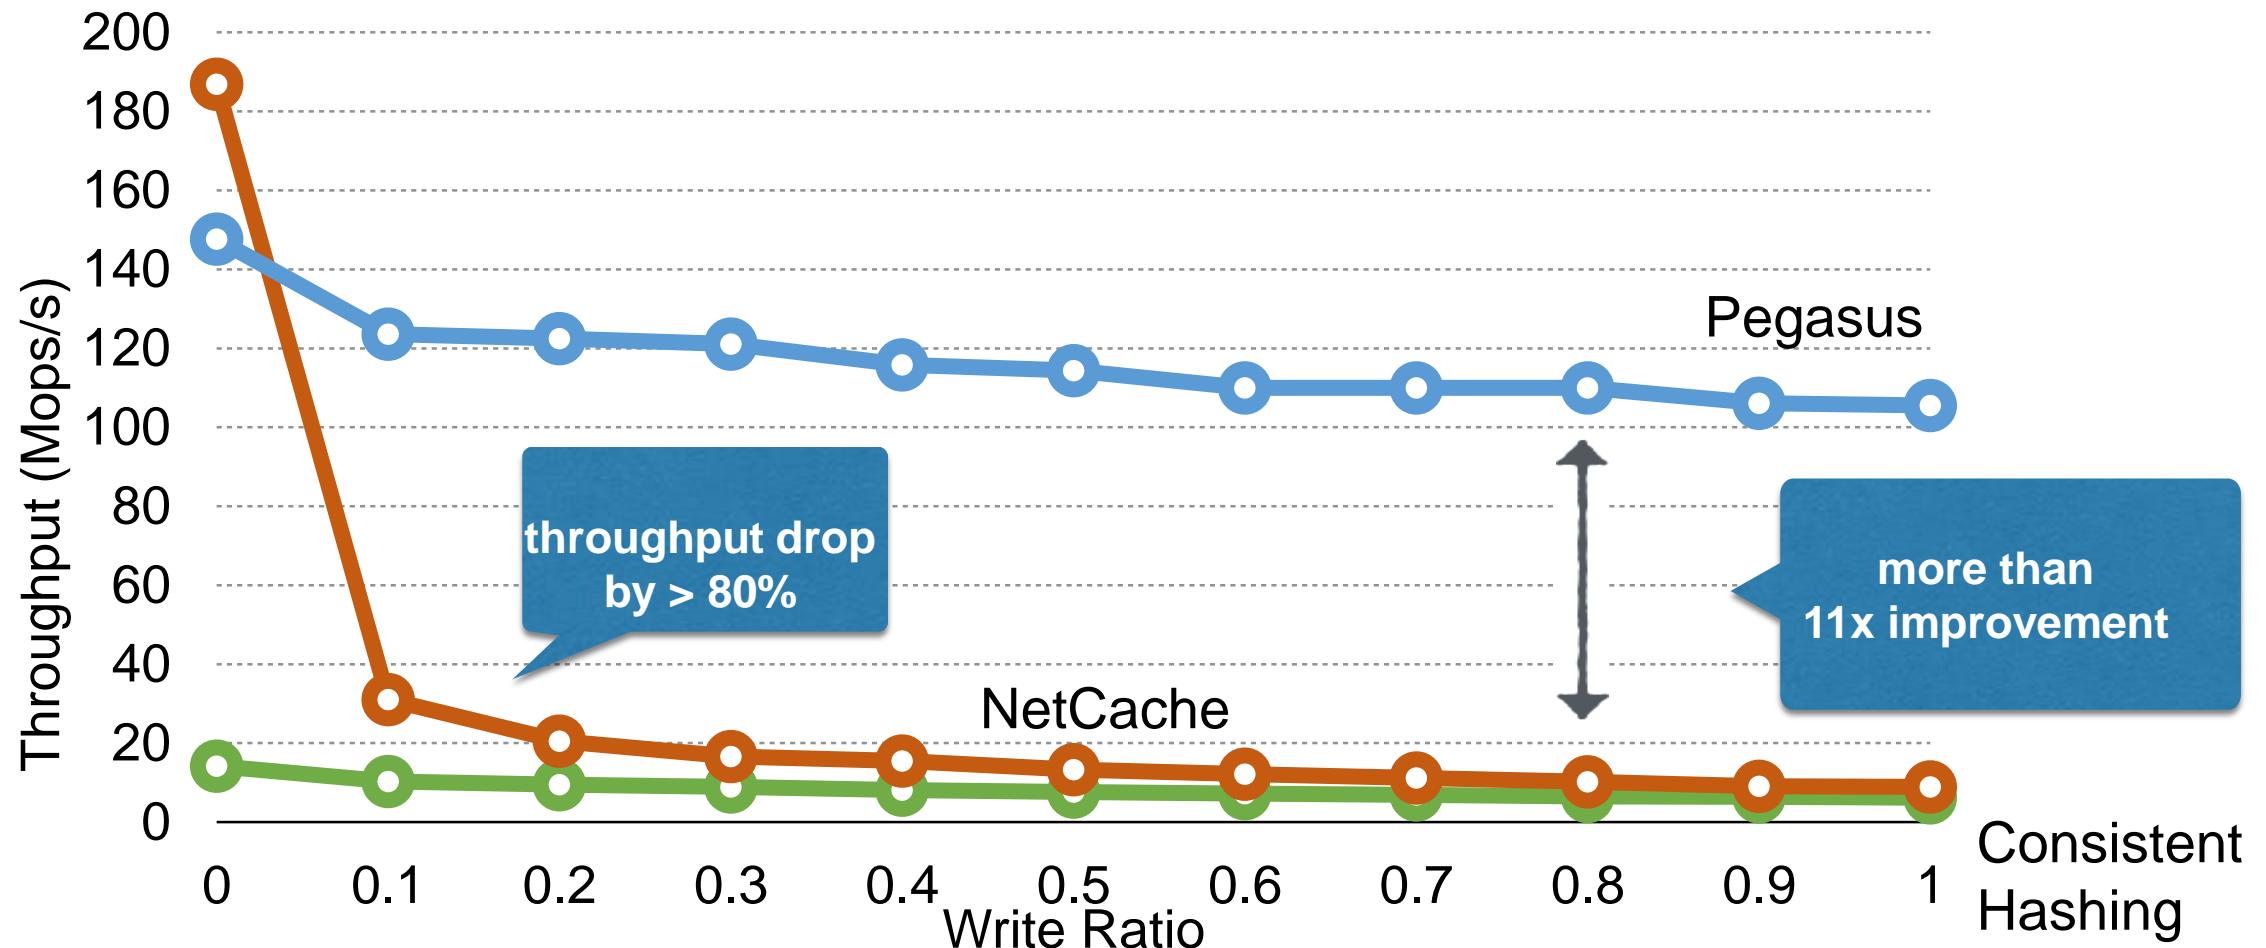
- Adding and removing replicated objects
  - Pegasus controller monitors object access frequencies
  - Updates coherence directory with most popular objects
- Server selection policy
  - Random
  - Weighted round-robin
- Handling server and rack failure
  - Multi-rack deployment
  - Each rack runs a Pegasus instance
  - Chain replication across racks

What benefits do we get?

# Pegasus is effective under highly skewed workloads



# Pegasus equally effective under different read/write ratios



# Conclusion

- Co-designing distributed systems with the network
  - Identify the right abstractions to be implemented **in the network**
  - Design protocols to take advantages of network primitives and provide **strong guarantees**
- Distributed systems with both **strong guarantees** and **high performance**
  - State machine replication: Network-Ordered Paxos
  - Distributed storage load balancing: Pegasus