

Fig 1.1. Elements of a modern computer system

(b)

OS

→ mapping function

• bidirectional (algm struct \leftrightarrow h/w)

• mapping efficiency

 → progr'ing

• mapping implies & includes

• processor scheduling

• memory maps

• interproc. comm.

} architectural
dependent
functions

Optimal mapping ?

Parallelism \rightarrow any stage. i.e.,

 design,
 prog., compile,
 run

Compiler support \rightarrow preprocessor

 → precompiler

 → parallelizing compiler

Table 1.1 Five Generations of Electronic Computers

Generation	Technology and Architecture	Software and Applications	Representative Systems
First (1945–54)	Vacuum tubes and relay memories, CPU driven by PC and accumulator, fixed-point arithmetic.	Machine/assembly languages, single user, no subroutine linkage, programmed I/O using CPU.	ENIAC, Princeton IAS, IBM 701.
Second (1955–64)	Discrete transistors and core memories, floating-point arithmetic, I/O processors, multiplexed memory access.	HLL used with compilers, subroutine libraries, batch processing monitor.	IBM 7090, CDC 1604, Univac LARC.
Third (1965–74)	Integrated circuits (SSI/-MSI), microprogramming, pipelining, cache, and lookahead processors.	Multiprogramming and time-sharing OS, multiuser applications.	IBM 360/370, CDC 6600, TI-ASC, PDP-8.
Fourth (1975–90)	LSI/VLSI and semiconductor memory, multiprocessors, vector supercomputers, multicomputers.	Multiprocessor OS, languages, compilers, and environments for parallel processing.	VAX 9000, Cray X-MP, IBM 3090, BBN TC2000.
Fifth (1991–present)	ULSI/VHSIC processors, memory, and switches, high-density packaging, scalable architectures.	Massively parallel processing, grand challenge applications, heterogeneous processing.	Fujitsu VPP500, Cray/MPP, TMC/CM-5, Intel Paragon.

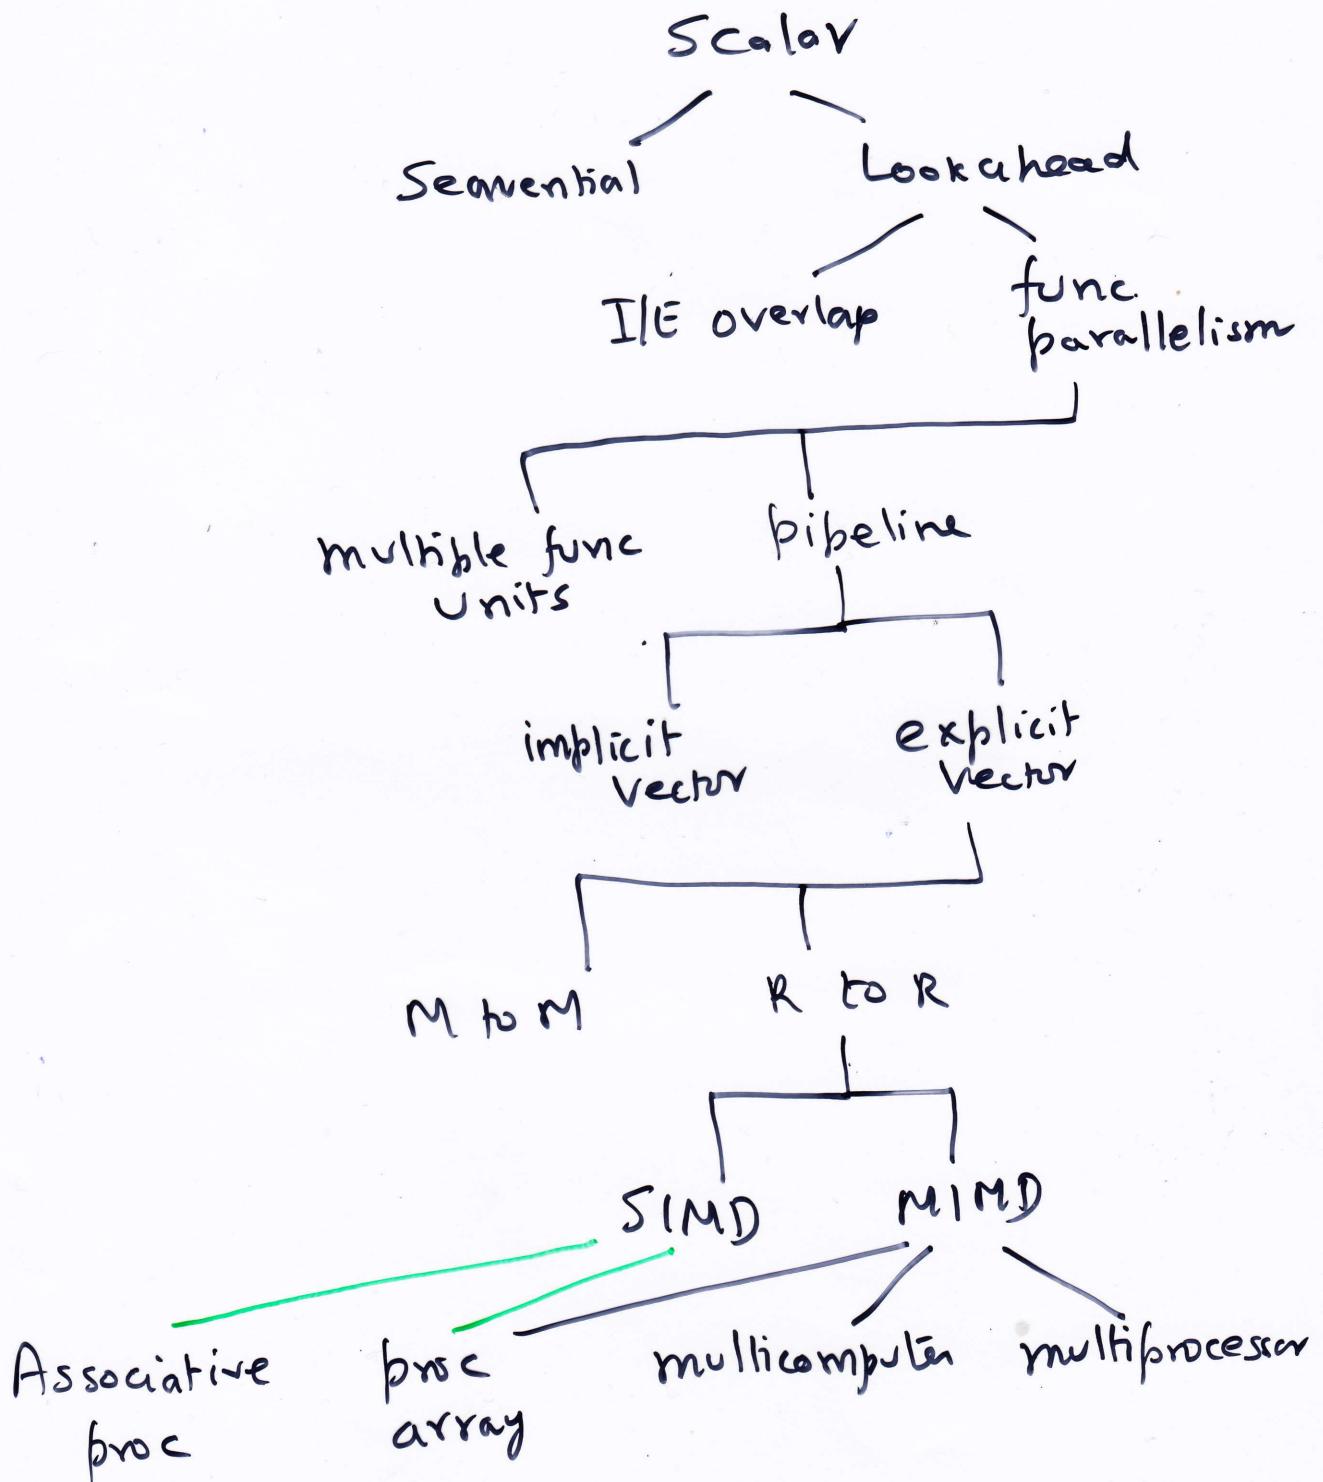


Table 1.2 Performance Factors Versus System Attributes

System Attributes	Performance Factors				Processor Cycle Time, τ
	Instr. Count, I_c	Average Cycles per Instruction, CPI	Processor Cycles per Instruction, τ	Memory References per Instruction, m	
Instruction-set Architecture	X	X			
Compiler Technology	X	X	X		
Processor Implementation and Control		X			X
Cache and Memory Hierarchy				X	X

Evolution of Computer Architecture

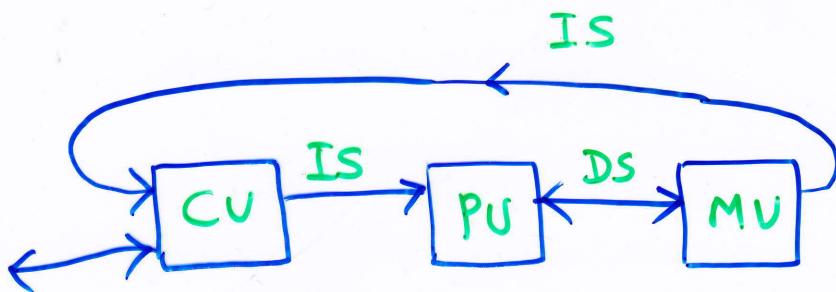
(fig 1.2, pg : 10)



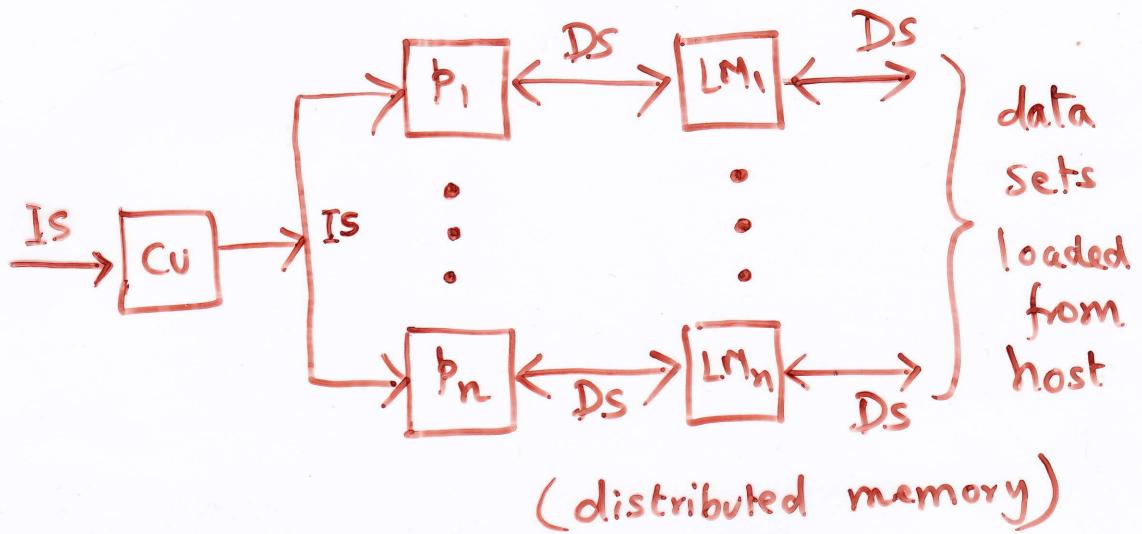
FLYNN's CLASSIFICATION

3

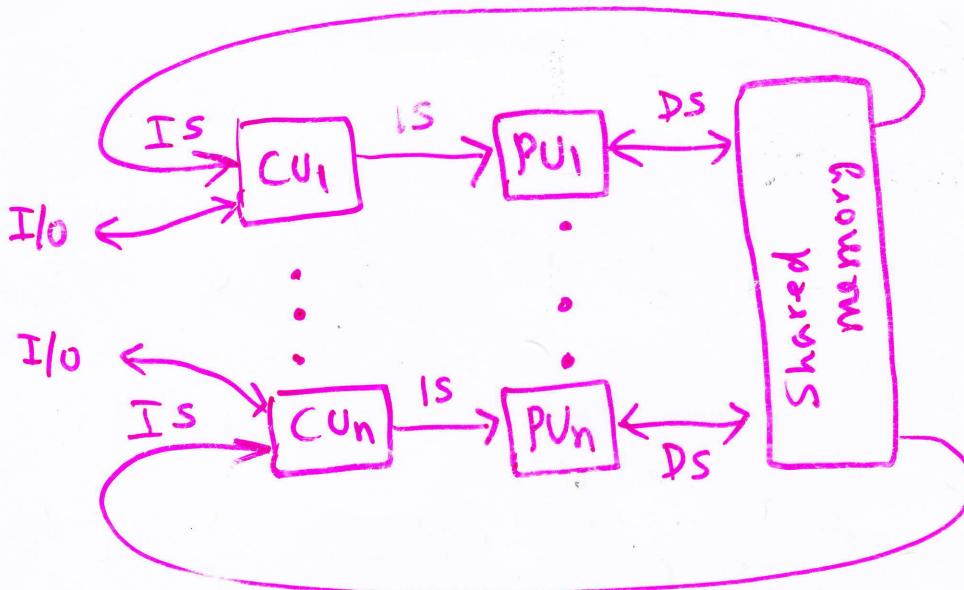
SISD



SIMD

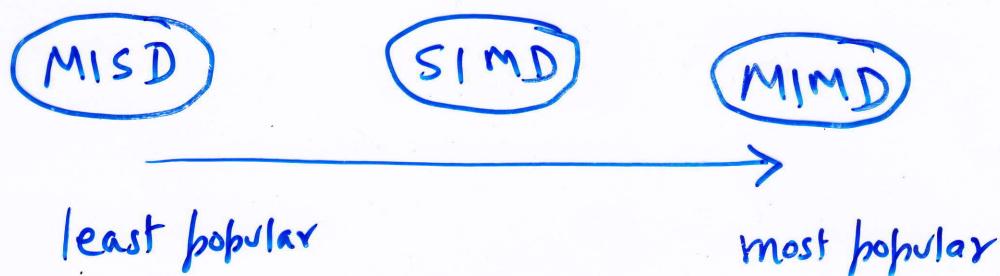


MIMD

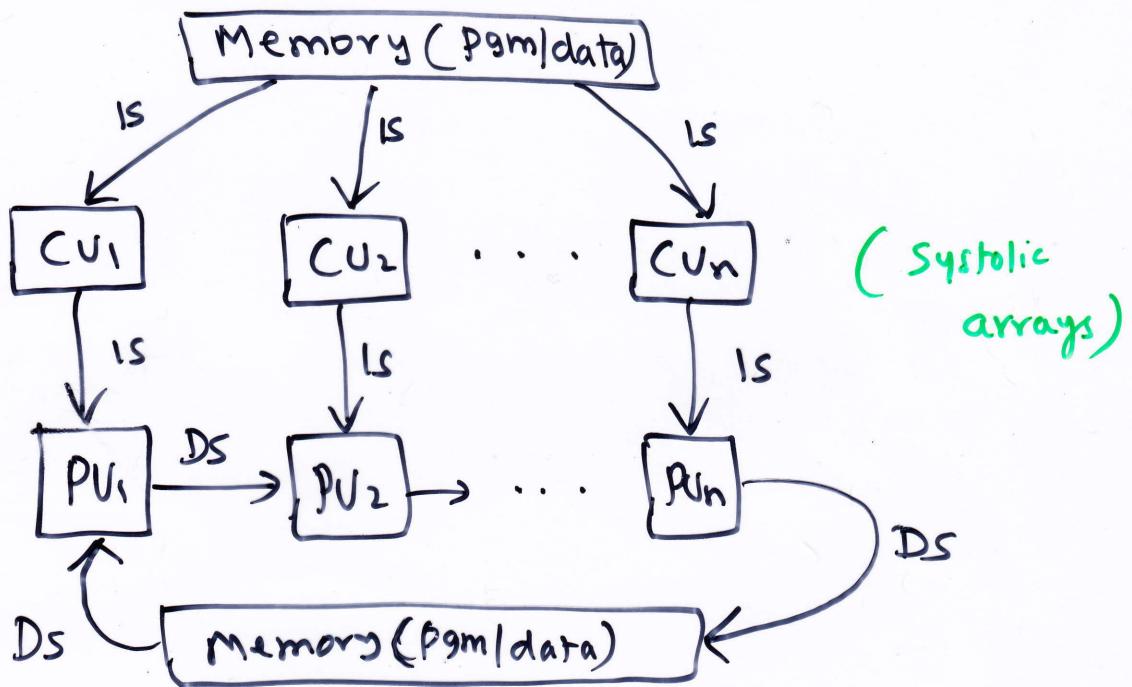


MIMD → general purpose computations

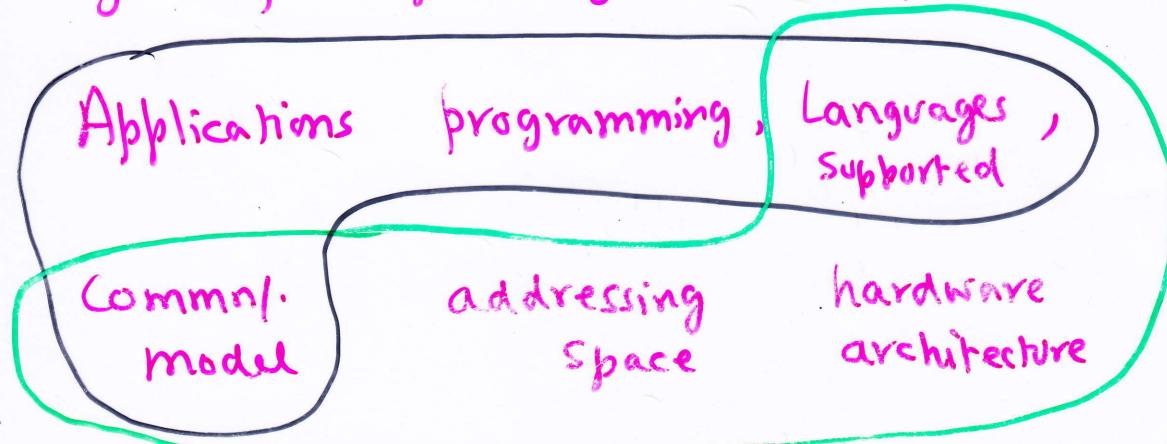
SIMD & MISD → special purpose



MISD:



Layers of computer system development



— : machine independent

— : machine dependent

Table 1.1 Five Generations of Electronic Computers

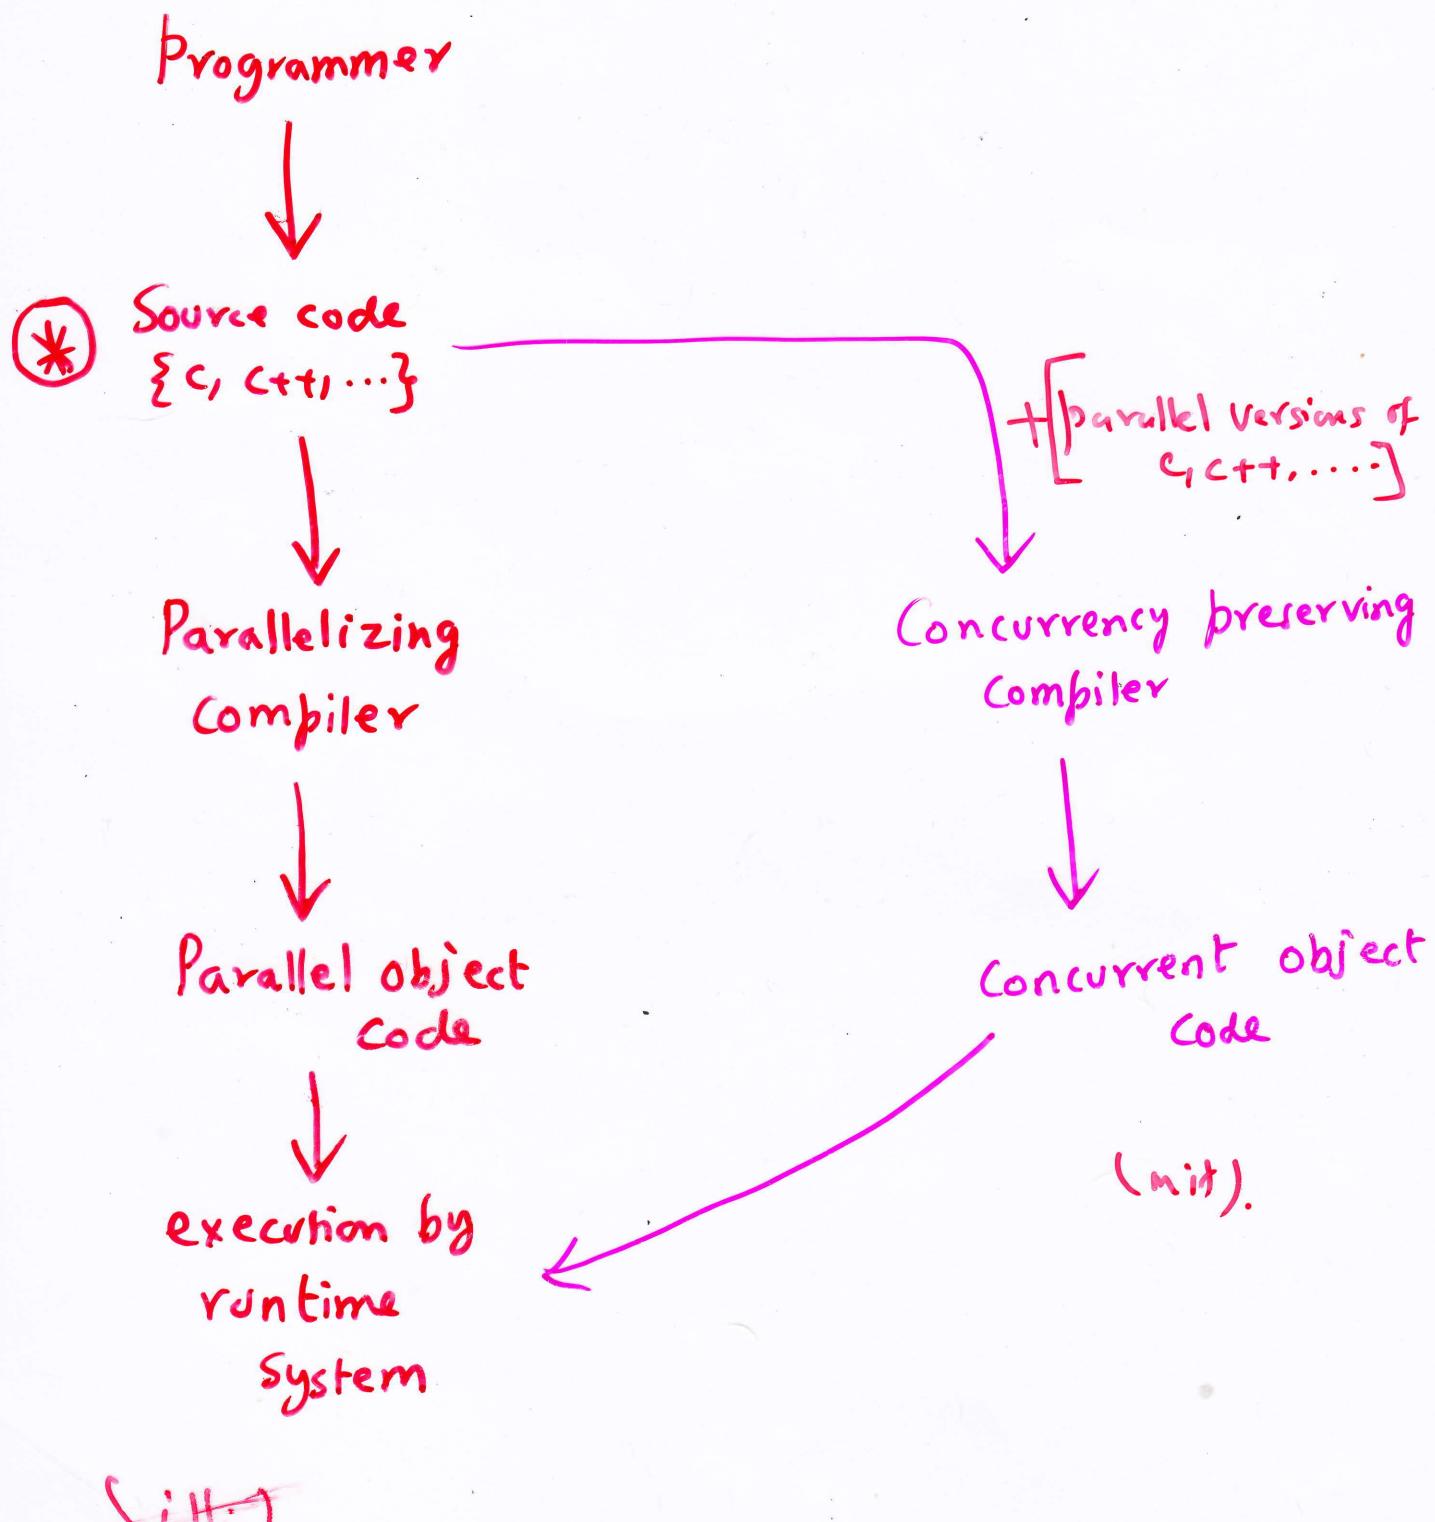
Generation	Technology and Architecture	Software and Applications	Representative Systems
First (1945–54)	Vacuum tubes and relay memories, CPU driven by PC and accumulator, fixed-point arithmetic.	Machine/assembly languages, single user, no subroutine linkage, programmed I/O using CPU.	ENIAC, Princeton IAS, IBM 701.
Second (1955–64)	Discrete transistors and core memories, floating-point arithmetic, I/O processors, multiplexed memory access.	HLL used with compilers, subroutine libraries, batch processing monitor.	IBM 7090, CDC 1604, Univac LARC.
Third (1965–74)	Integrated circuits (SSI/-MSI), microprogramming, pipelining, cache, and lookahead processors.	Multiprogramming and time-sharing OS, multiuser applications.	IBM 360/370, CDC 6600, TI-ASC, PDP-8.
Fourth (1975–90)	LSI/VLSI and semiconductor memory, multiprocessors, vector supercomputers, multicomputers.	Multiprocessor OS, languages, compilers, and environments for parallel processing.	VAX 9000, Cray X-MP, IBM 3090, BBN TC2000.
Fifth (1991–present)	ULSI/VHSIC processors, memory, and switches, high-density packaging, scalable architectures.	Massively parallel processing, grand challenge applications, heterogeneous processing.	Fujitsu VPP500, Cray/MPP, TMC/CM-5, Intel Paragon.



Table 1.2 Performance Factors Versus System Attributes

System Attributes	Performance Factors				Processor Cycle Time, τ
	Average Cycles per Instruction, CPI				
	Processor Cycles per Instruction, τ	Memory References per Instruction, m	Memory Access Latency, k		
Instruction-set Architecture	X	X			
Compiler Technology	X	X	X		
Processor Implementation and Control		X			X
Cache and Memory Hierarchy				X	X

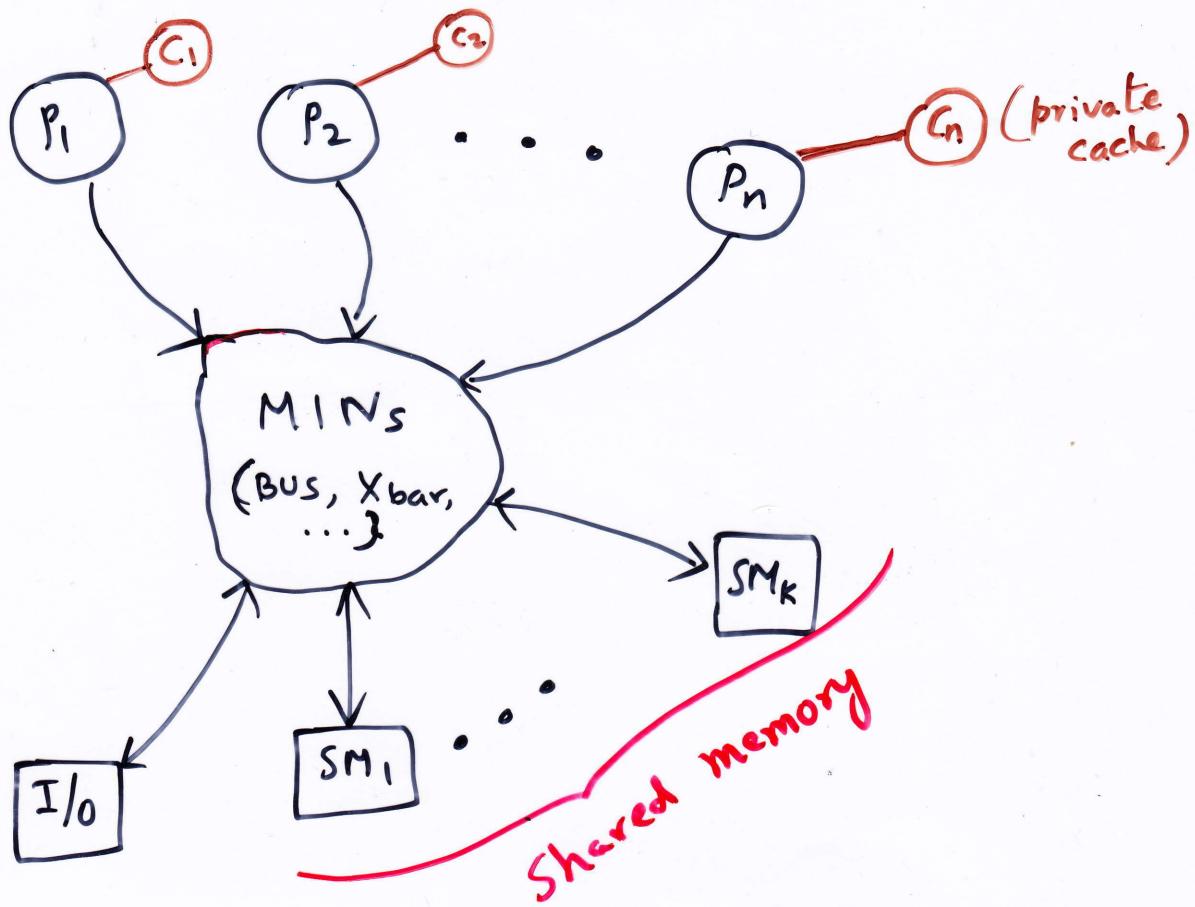
Two Approaches to parallel programming



(*) Implicit parallelism → Sequential languages

Explicit parallelism → Concurrent dialects
of C, C++, Pascal...

UMA Multiprocessor Model



Example (1.2)

1. DO 10 I=1,N

2. A(I) = B(I) + C(I) •

3. 10 CONTINUE

4. SUM = 0 •

5. DO 20 J=1,N

6. Sum = Sum + A(J) •

7. 20 CONTINUE

Doall k=1,M

Do 10 I=L(k-1)+1, KL

A(I) = B(I) + C(I)

10 CONTINUE

SUM(K) = 0

DO 20 J=1,L

SUM(K) = SUM(K) +
A(L(K-1)+J)

20 CONTINUE

Endall

$$(L = N/M)$$

Sequential processing: $2N$ cycles (8)

($L_2, L_4, L_6 \rightarrow 1$ m/c cycle to execute

$L_1, L_3, L_5, L_7 \rightarrow$ ignored for analysis)

Multiprocessing

- Divide the looping into M sections : $L = (N/M)$
- Doall declares that all M sections are executed by M processors in parallel

I loop : L cycles

J loop : M partial sums in L cycles

$2L$ cycles are consumed to produce M partial sums. We need to merge these M to get the final sum.

$(N=2^{20})$

$\cdot l$ -level binary adder : $\log_2 M = l$

each pair : k cycles $\Rightarrow (k+1)l$ cycles

$$\Rightarrow \textcircled{2L} + (k+1)l = \left(2 \frac{N}{M} + (k+1) \log_2 M \right) \quad //$$

(9)

$$\text{Speed-up} = \frac{2N}{2 \cdot L + (k+1) \log_2 M}$$

where $L = (N/M)$

$$\underline{N=2^{20}} \Rightarrow$$

($k=200$ cycles,
 $M=256$)

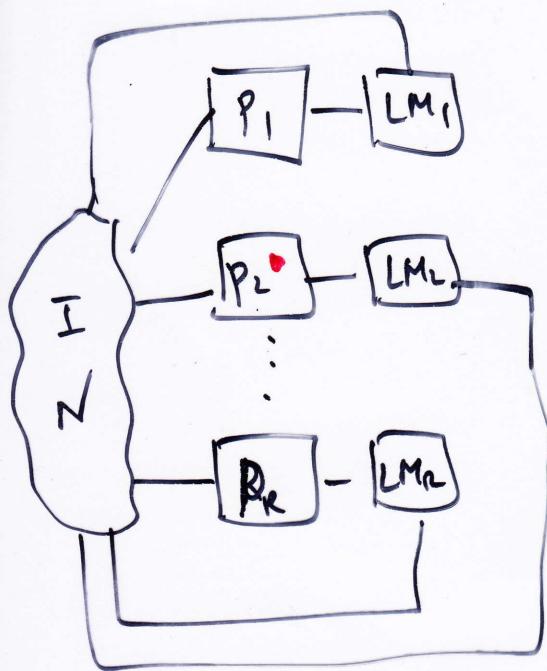
$$\text{Speed up} = \frac{2^{21}}{(2^{13} + 1608)}$$

$$\approx \underline{\underline{83.6\%}}$$

NUMA Model

Shared mem System → access times vary with the location of a word.

(1)



- Collection of all local memories forms a global address space accessible by all the processors

$$t(LM) < t(GM)$$

→ true with
the current day
R.C architectures!

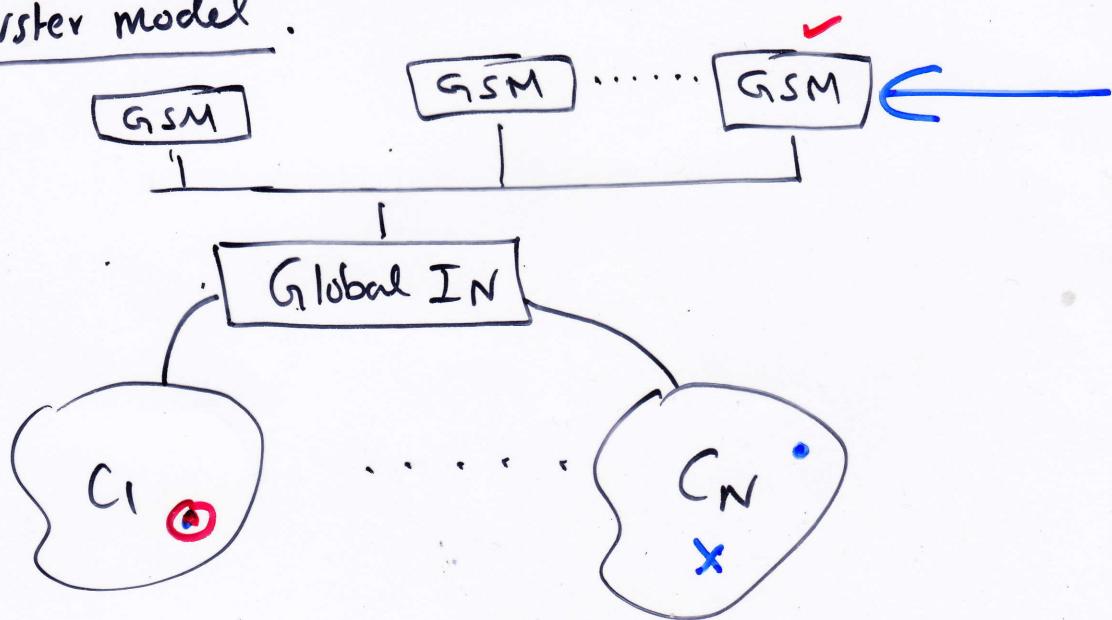
modern day → RCA

$$t(LM) < t(GM) < t(\text{disk})$$

Shared local mem. model

(2)

Cluster model:



Each $C_i \rightarrow$ Shared mem model or $C_i \rightarrow$ can be a VMA

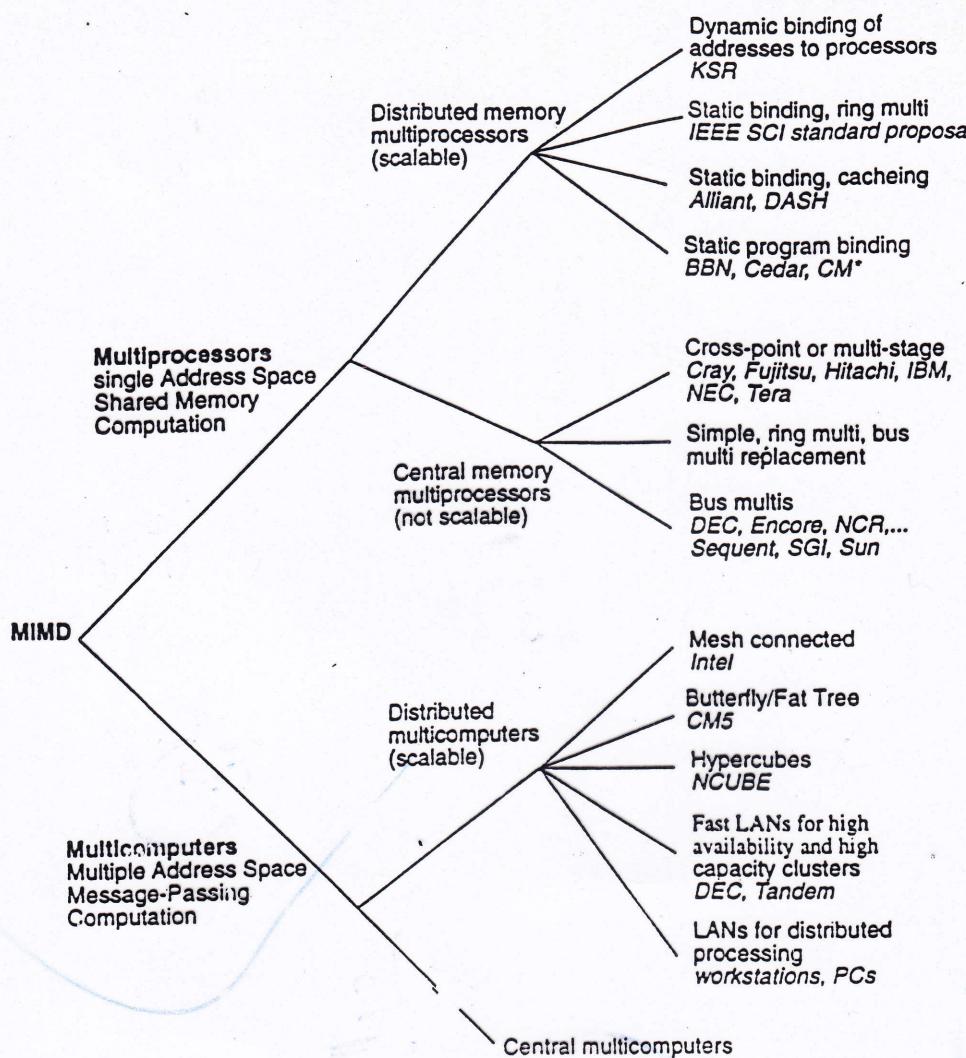


Figure 1.10 Bell's taxonomy of MIMD computers. (Courtesy of Gordon Bell; reprinted with permission from the *Communications of ACM*, August 1992)

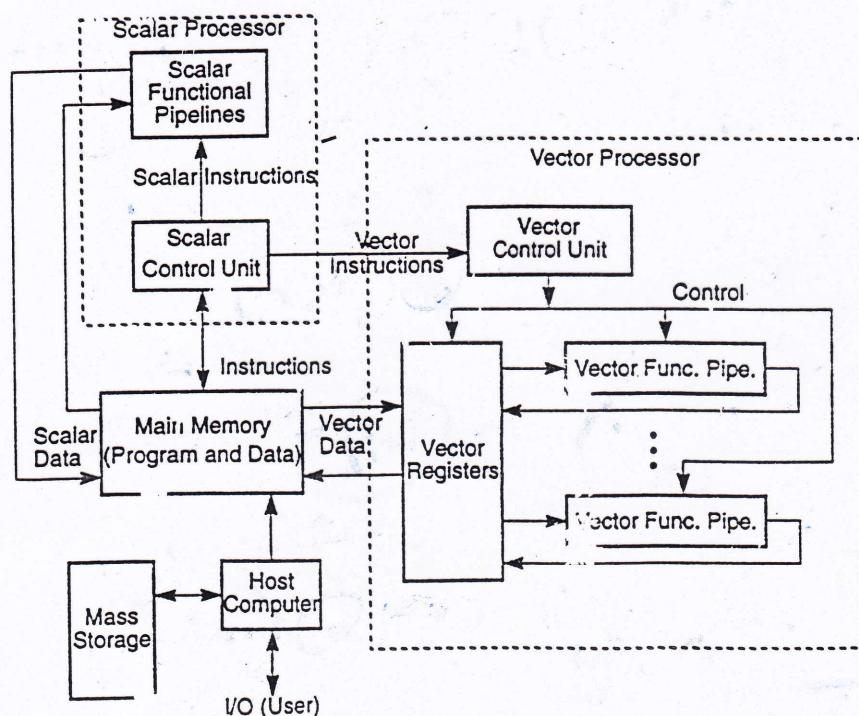
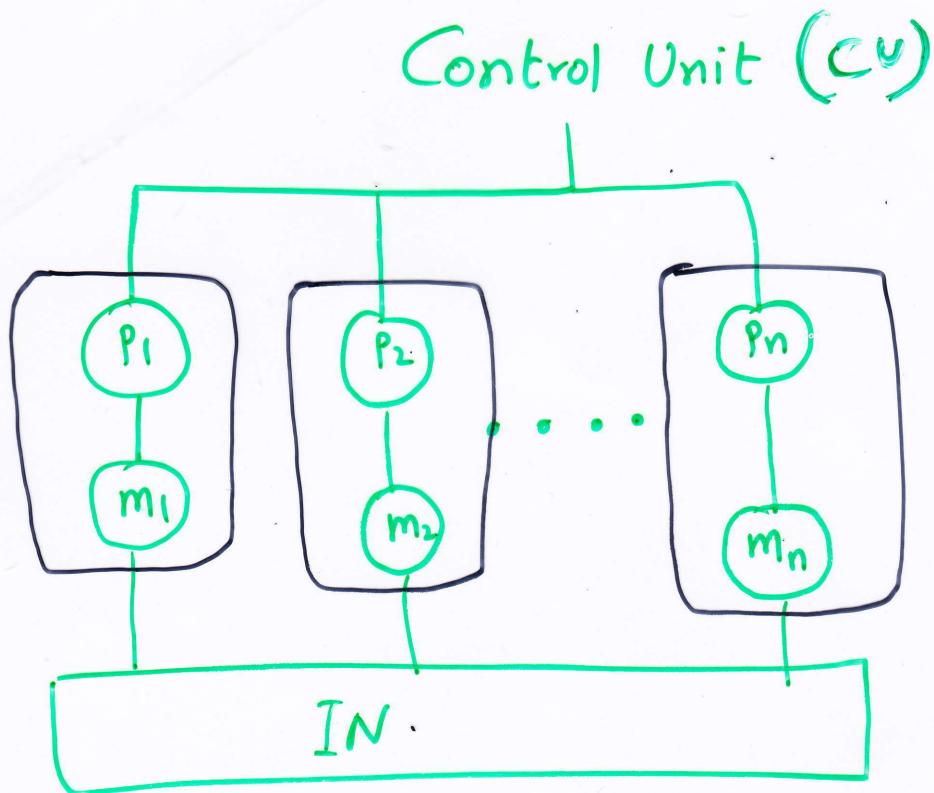


Figure 1.11 The architecture of a vector supercomputer.

SIMD Machine Model



$$M = \langle N, C, I, M, R \rangle$$

N : # of processing elements

C : Instr. directly executed by CU

I : Instr. broadcast by CU to all PEs.

M : Masking schemes

R : data-routing functions